

COMMONWEALTH OF AUSTRALIA
Copyright Regulation 1969

WARNING

This material has been copied and communicated to you by or on behalf
of **Curtin University of Technology** pursuant to Part VB of the
Copyright Act 1968 (**the Act**)

The material in this communication may be subject to copyright under the
Act. Any further copying or communication of this material by you
may be the subject of copyright protection under the Act.

Do not remove this notice

Operating Systems

COMP2006

CPU Scheduling

Lecture 4

CPU Scheduling

References: Silberschatz, Galvin, and Gagne, *Operating System Concepts*, Chapter 6

Topics:

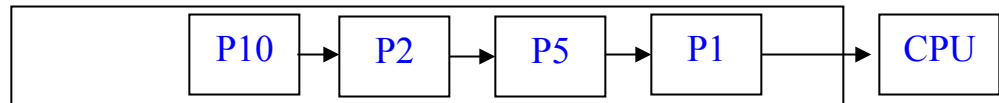
- ★ Scheduling concepts.
- ★ CPU scheduling algorithms.
- ★ CPU scheduling evaluations.

CPU Scheduling

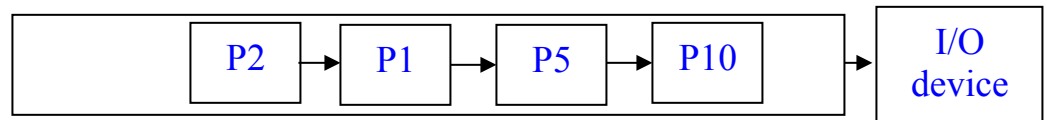
Why schedule the CPU?

- ★ The objective of multiprogramming is to maximize CPU utilization.
 - Select and run one process in ready queue when CPU is available.
- ★ Process execution consists of a cycle of CPU execution and I/O wait; *i.e.*, CPU and I/O burst cycle.
 - Processes alternate between these two activities
 - A process terminates after the last CPU burst
- ★ Processes can be:
 - CPU-bound process: very long CPU burst.
 - I/O-bound process: short CPU burst.

Ready Queue:



Device Queue:



Why (cont.)

Alternating sequence of CPU and I/O bursts

```
x := 0;
read from file
```

Wait for I/O

```
x := x + y;  
write to file
```

Wait for I/O

```
x := x + z;  
write to file
```

CPU burst

I/O burst

CPU burst

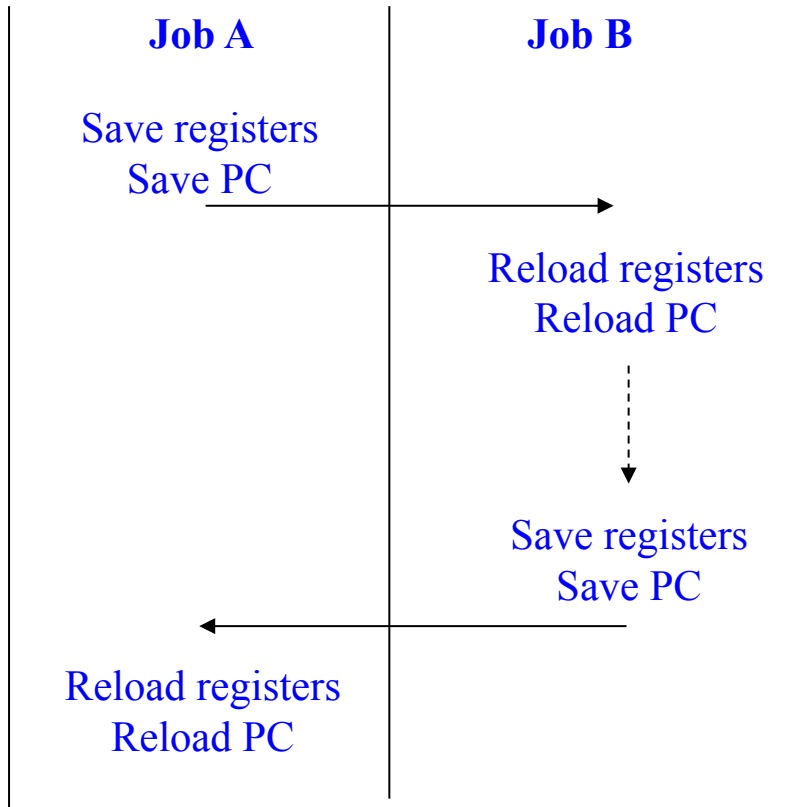
I/O burst

CPU burst

Histogram of CPU-burst Times



How and when do we switch CPU?

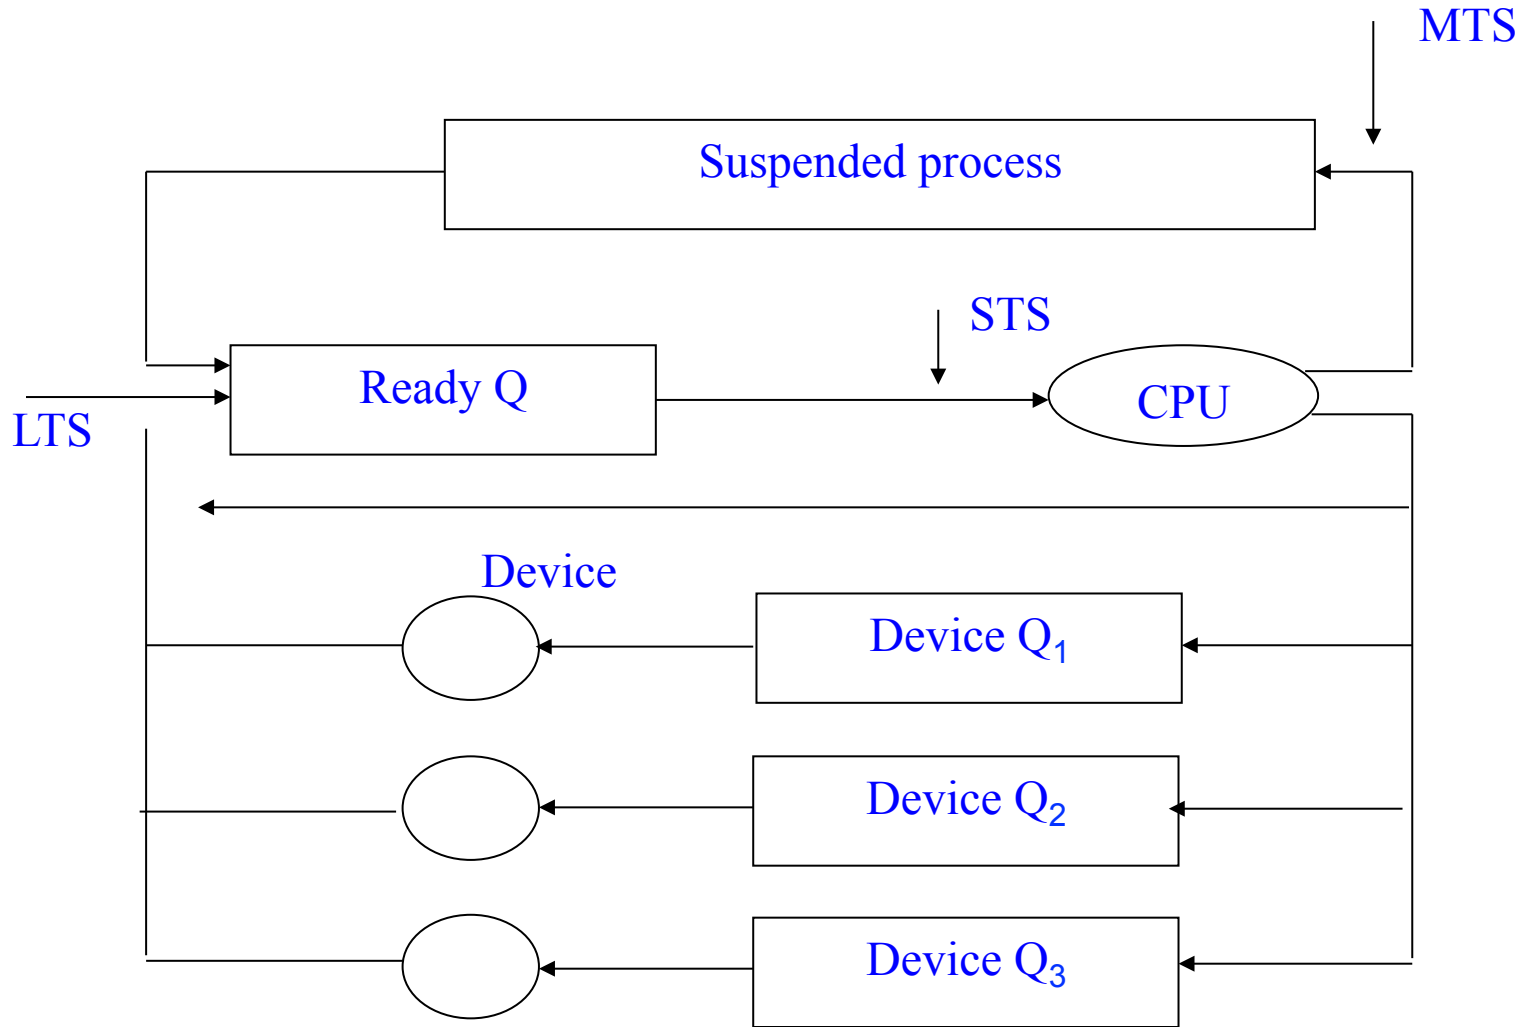


- From running to waiting state (*e.g.*, I/O request, or calling wait()).
- From running to ready state (*e.g.*, timer off).
- Process terminates.

CPU schedulers

- ★ CPU scheduler selects one of the processes in memory ready for execution, and allocates the CPU to it.
 - also called as Short Term Scheduler (STS)
 - Ready queue can be a FIFO, priority, a tree, or unordered linked list.
 - The content of the queue: PCBs of the processes.
 - This scheduler must be quick. Why?
- ★ Other schedulers:
 - Medium term scheduler (MTS) → Swaps jobs in and out of memory to reduce contention for the CPU.
 - Long term scheduler (LTS) → Determines which jobs are admitted
 - LTS is executed less frequently than STS
 - It is invoked only when a job finishes.
 - LTS controls the degree of multiprogramming
 - It must select carefully between CPU bound and I/O bound jobs.
 - LTS may be absent on time-sharing systems
 - MTS is added in this case.

Queueing diagram



Preemption

- ★ Scheduling can be *preemptive* or *non-preemptive*.
- ★ Non-preemptive: once the CPU has been allocated to a process, the process keeps the CPU until it releases the CPU either by terminating or by switching to the waiting state.
 - Example: Windows 3.x
- ★ Preemptive: otherwise.
 - Example: Windows 95 onward, Mac OS X, Linux.
- ★ Preemptive scheduling needs special hardware (*e.g.*, timer).
- ★ Preemptive scheduling can result in a race condition.
 - What happens when one process updating a shared data is preempted, and the second process is accessing the data?
- ★ Be careful when pre-empting a kernel process
 - What happens if the kernel is in the middle of changing important kernel data involving the process? See page 186.

Dispatcher

- ★ Dispatcher is the module that gives control of the CPU to the process selected by the short-term scheduler.
- ★ This function involves:
 - Switching context.
 - Switching to user mode.
 - Jumping to proper location in the user program to restart that program.
- ★ Dispatch latency: the time it takes for the dispatcher to stop one process and start running another
 - must be short because every context switch invokes dispatcher.

Scheduling Criteria

- ★ Criteria are used for comparing the scheduling algorithms to determine which algorithm is better.
- ★ Some scheduling criteria include:
 - **CPU utilisation:** percent usage of CPU → the higher the better.
 - **Throughput:** #processes that complete their execution per time unit → the higher the better.
 - **Turnaround time:** amount of time to execute a particular process, *i.e.*, sum of the periods spent waiting to get into memory, waiting in the ready queue, executing on the CPU, and doing I/O → the shorter the better.
 - **Waiting time:** the sum of the periods a process spent waiting in the ready queue → the shorter the better.
 - **Response time:** amount of time it takes from when a request was submitted until the first response is produced, NOT output → for time sharing environment.

Criteria (cont.)

- ★ For interactive systems:
 - Minimise variance in response time.
 - ★ A system with predictable response is more useful than a system that is faster on average.
 - In general: minimise the maximum response time *or* average response time.

- ★ Should we also include *fairness* in the criteria?
 - Make sure that each process gets its fair share of the CPU time.

Scheduling Algorithms

(a) First Come First Served (FCFS) Scheduling

- A process that requests CPU first is allocated the CPU first.
- Easy implementation with a FIFO queue.
- Non-preemptive.
- Performance: poor waiting time, turnaround time and response time.

Process	Burst time
1	24
2	3
3	3

Examples

Case (i): Processes arrive in the order P1, P2, P3

The Gantt chart for the schedule is:



★ Waiting time for P₁ = 0; P₂ = 24; P₃ = 27.

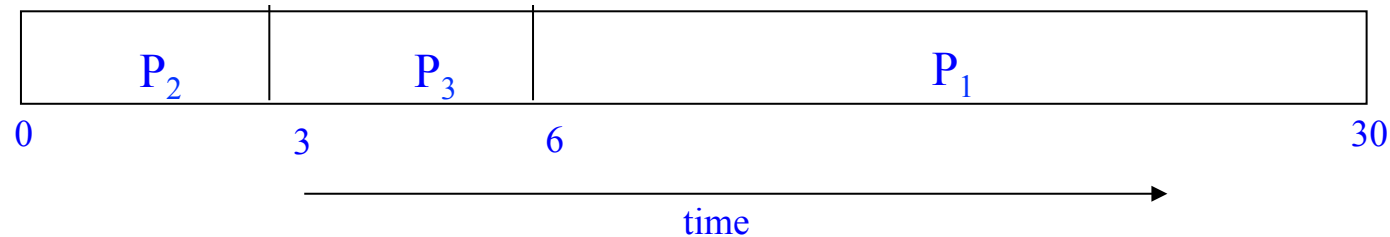
★ Average waiting time: $\frac{0 + 24 + 27}{3} = \frac{51}{3} = 17$

★ Average turnaround time: $\frac{24 + 27 + 30}{3} = \frac{81}{3} = 27$

★ Waiting time = turnaround time – burst time

Examples (cont.)

Case (ii) Processes arrive in the order P2, P3, P1



Waiting time for P1 = 6; P2 = 0; P3 = 3.

$$\text{Average waiting time} = \frac{6 + 0 + 3}{3} = \frac{9}{3} = 3 \rightarrow \text{better than (i).}$$

$$\text{Average turnaround time: } \frac{3 + 6 + 30}{3} = \frac{39}{3} = 13$$

Convoy effect: short processes behind long process.

Consider one CPU bound job and several I/O bound jobs \rightarrow lower CPU utilisation and device utilisation.

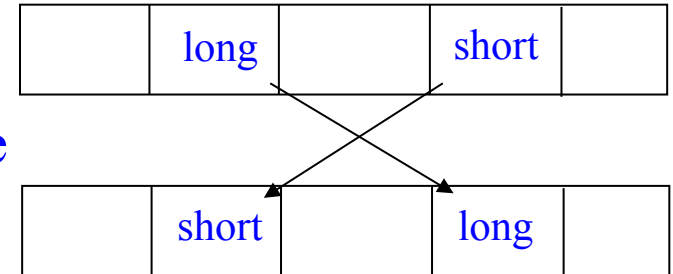
Algorithms (cont.)

(b) Shortest Job First (SJF) Scheduling

- Associate with each process the length of its next CPU burst, and use this length to schedule the process with shortest time.
- Two schemes:
 - ★ Non-preemptive – once the CPU is given to the process, it cannot be pre-empted until it completes its CPU burst.
 - ★ Preemptive – if a new process arrives with CPU burst length less than remaining time of current executing process, preempt
 - Also known as the Shortest-Remaining-Time-First (SRTF).
- SJF is optimal – gives minimum average waiting time, turnaround time, and response time **for a given set of processes.**

SJF (cont.)

- ★ Moving a short job before a long one decreases the waiting time of the short job more than it increases the waiting time of the long job → therefore the average waiting time decreases.



Simple proof

Consider the case of four processes with run times of a , b , c , and d . The first process finishes at time a , the second process finishes at time $a + b$, etc. The average turn around time is
$$\frac{4a + 3b + 2c + d}{4}.$$

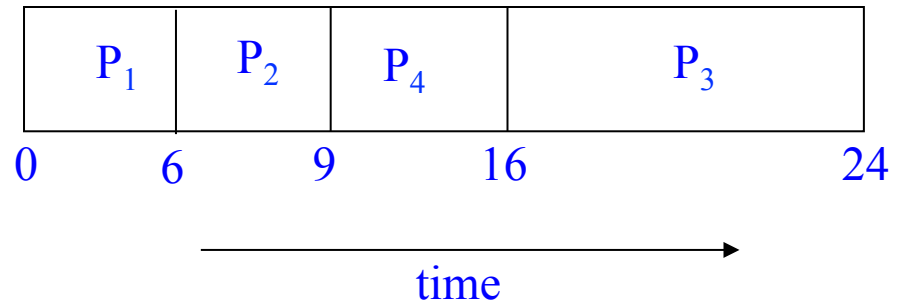
It is clear that a contributes more to the average than the other times, so it should be the shortest process, with b next, then c , and finally d as the longest since it affects only its own turn around time.

Problem: How to know which of the currently run-able processes have the shortest CPU burst?

Example (non-preemptive SJF)

Process	Arrival time	Burst time
1	0	6
2	1	3
3	2	8
4	3	7

The Gantt chart for the schedule:



Waiting time for $P_1 = 0$; $P_2 = 6 - 1$; $P_3 = 16 - 2$; $P_4 = 9 - 3$.

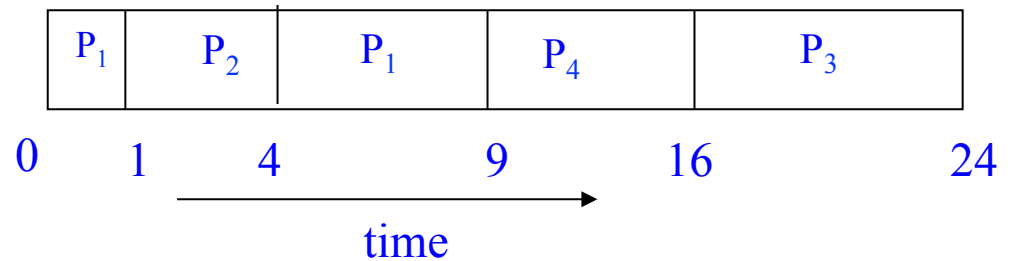
$$\text{Average waiting time} = \frac{0 + 5 + 14 + 6}{4} = \frac{25}{4} = 6.25$$

$$\text{Average turnaround time: } \frac{(6 - 0) + (9 - 1) + (16 - 3) + (24 - 2)}{4} = \frac{49}{4} = 12.25$$

Example (preemptive SJF)

Process	Arrival time	Burst time
1	0	6
2	1	3
3	2	8
4	3	7

The Gantt chart for the schedule:



Waiting time for P1 = 4-1; P2 = 1-1; P3 = 16-2; P4 = 9-3

$$\text{Average waiting time} = \frac{3 + 0 + 14 + 6}{4} = \frac{23}{4} = 5.75$$

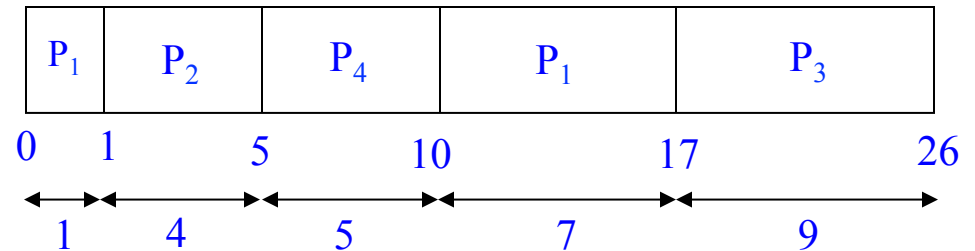
$$\text{Average turnaround time: } \frac{(9 - 0) + (4 - 1) + (24 - 2) + (16 - 3)}{4} = \frac{47}{4} = 11.75$$

Note: waiting time of a process is its turnaround time *minus* its burst time.

Other example (SJF – preemptive)

Process	Arrival time	Burst time
1	0	8
2	1	4
3	2	9
4	3	5

The Gantt chart for the schedule:



Average turnaround time:
$$\frac{(17 - 0) + (5 - 1) + (26 - 2) + (10 - 3)}{4} = 13$$

Average turnaround time for non-pre-emptive SJF: 14.25 seconds.

Determining Length of next CPU Burst

Problems: How do we know the length of the job?

- ★ Can only estimate the length
- ★ Can be done using the length of previous CPU burst, using exponential averaging

$$\tau_{n+1} = \alpha.t_n + (1 - \alpha).\tau_n$$

t_n = actual length of n^{th} CPU burst.

τ_{n+1} = predicted value for the next burst.

$\alpha : 0 \leq \alpha \leq 1$ controls the weight of the recent vs. past history.

- ★ More commonly $\alpha = 1/2$, so recent history and past history are equally weighted; The initial τ_0 can be defined as a constant or as an overall system average.

Examples of Exponential Averaging

$\alpha = 0 \rightarrow \tau_{n+1} = \tau_n$; recent history does not count

$\alpha = 1 \rightarrow \tau_{n+1} = t_n$; only the actual last CPU burst counts

If we expand the formula, by repeatedly substituting $\tau_n = \alpha t_{n-1} + (1 - \alpha) \tau_{n-1}$,

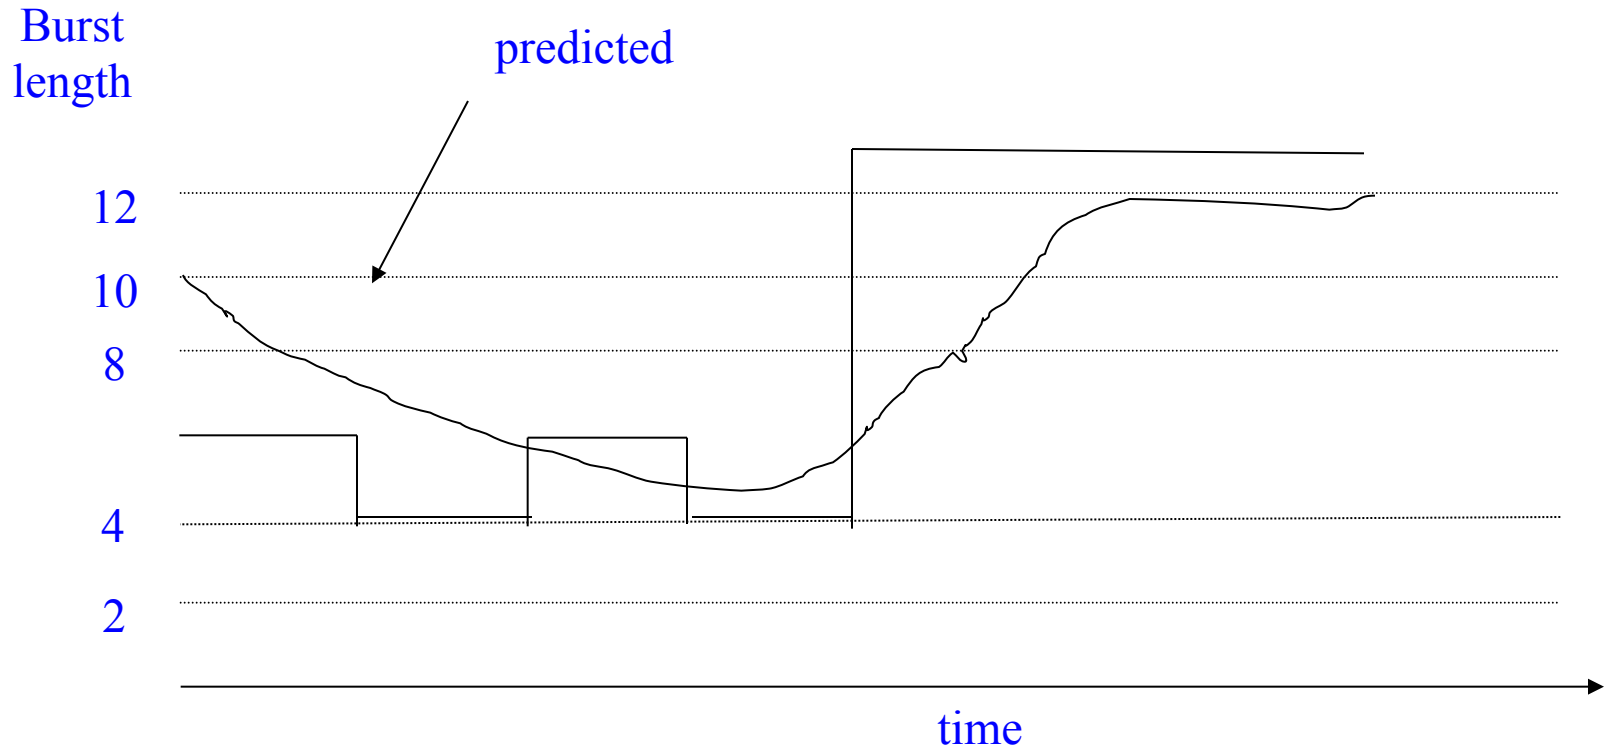
we get:

$$\tau_{n+1} = \alpha t_n + (1 - \alpha) \alpha t_{n-1} + \dots + (1 - \alpha)^j \alpha t_{n-j} + \dots + (1 - \alpha)^{n+1} \tau_0$$

Since both α and $(1 - \alpha)$ are less than or equal to 1, each successive term has less weight than its predecessor.

Example: $\alpha = \frac{1}{2}$ and $\tau_0 = 10$

CPU		6	4	6	4	13	13	13
Guess	10	8	6	6	5	9	11	12



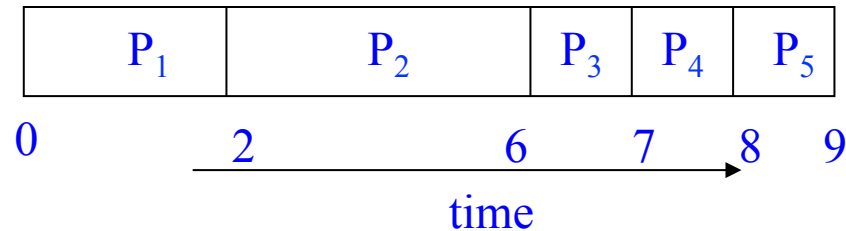
SJF (cont.)

- ★ SJF is only optimal when all jobs are available simultaneously.

Process	Arrival time	Burst time
1	0	2
2	0	4
3	3	1
4	3	1
5	3	1

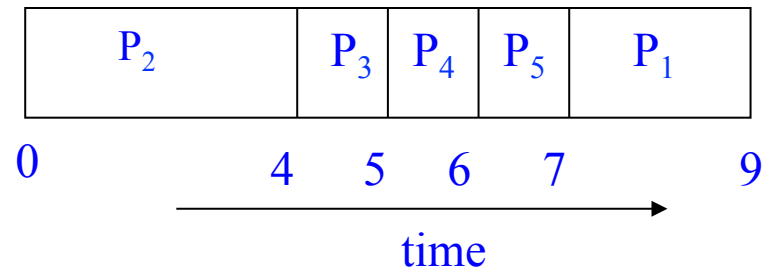
If we run them in the order 1, 2, 3, 4, 5,
average waiting time =

$$\frac{(0) + (2 - 0) + (6 - 3) + (7 - 3) + (8 - 3)}{5} = 2.8$$



If we run them in the order 2, 3, 4, 5,
1 (*non- SJF*), we get a *better* average
waiting time =

$$\frac{(0) + (4 - 3) + (5 - 3) + (6 - 3) + (7 - 0)}{5} = 2.6$$



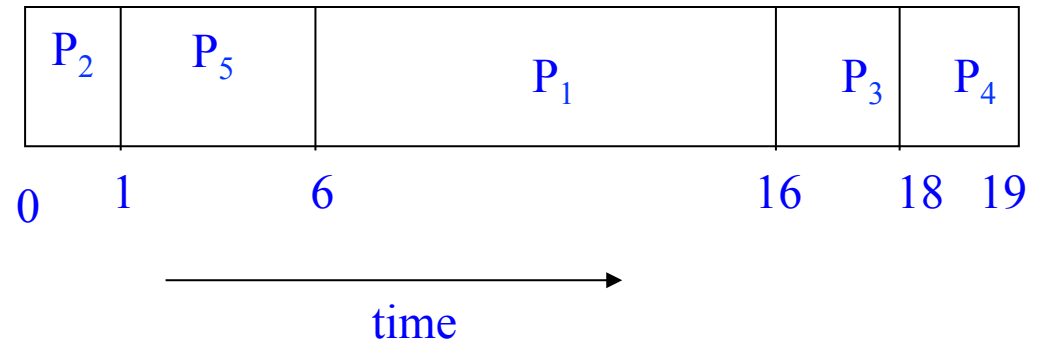
Priority Scheduling

- ★ A priority number (integer) is associated with each process.
- ★ The CPU is allocated to the process with the highest priority (low number → high priority).
 - Preemptive → preempt the CPU if the priority of the arrived process is higher than the priority of the currently running process.
 - Non-preemptive.
- ★ Internally defined priority - based on some measurable quantity:
 - Time limits, memory requirements, number of open files, ratio of average I/O to CPU bursts.
- ★ Externally defined priority:
 - Type and amount of funds, department, politics.
- ★ Problem: starvation → low priority processes may never execute.
 - Solution: aging → as time progresses increase the priority of the process.
- ★ Shortest Job First is a priority scheduling
 - priority is the predicted next CPU burst time.

Example (non-preemptive priority scheduling)

Process	Burst time	Priority
1	10	3
2	1	1
3	2	3
4	1	4
5	5	2

The Gantt chart for the schedule



Waiting time for P₁ = 6; P₂ = 0; P₃ = 16; P₄ = 18; P₅ = 1.

$$\text{Average waiting time} = \frac{6 + 0 + 16 + 18 + 1}{5} = \frac{41}{5} = 8.2$$

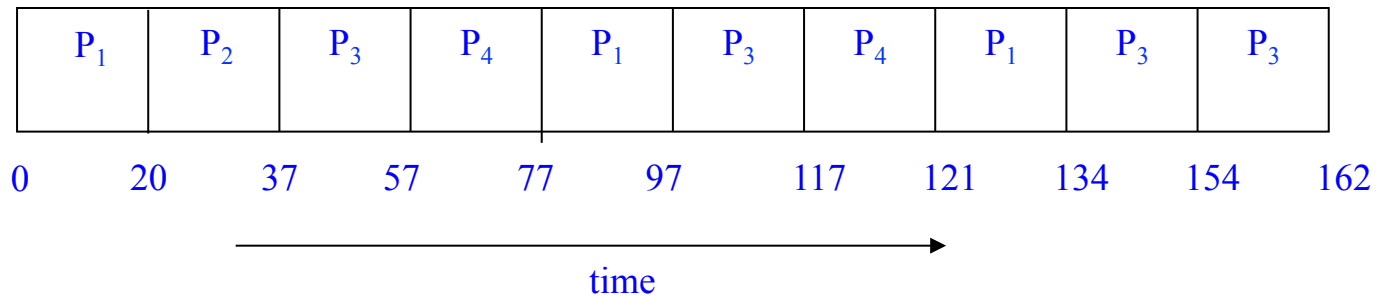
Round Robin (RR) Scheduling

- ★ Each process gets a small unit of CPU time (time quantum), usually 10 to 100 ms.
 - After this time has elapsed, the process is pre-empted and added to the end of the ready queue.
- ★ If the ready queue has n processes and the time quantum is q ,
 - each process gets $1/n$ of the CPU time in chunks of at most q time units at once.
 - No process waits more than $(n-1) q$ time unit.
- ★ Performance of RR system depends on the size of q .
 - Large $q \rightarrow$ FCFS.
 - Small $q \rightarrow q$ must be large with respect to context switch,
 - ★ Otherwise overhead is too high.
 - ★ In practice context switch is less than $10\mu\text{s}$.
 - Turn around time improves if most processes finish their next CPU burst within q .
 - Rule of thumb: 80% of CPU burst is at most q .
- ★ Typically, higher average turnaround than SRTF, but better response time.

Example (time quantum=20)

Process	Burst time
P ₁	53
P ₂	17
P ₃	68
P ₄	24

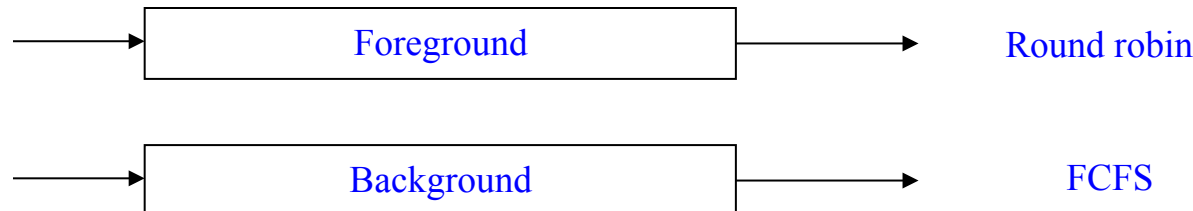
The Gantt chart for the schedule is:



Multilevel Queue Scheduling

- ★ Ready queue is partitioned into separate queues; e.g., foreground (interactive), background (batch).
- ★ Each queue has its own scheduling algorithm.
- ★ Scheduling must be done between the queues.
 - Fixed priority scheduling; *i.e.*, serve all from foreground then from background. Possibility of starvation.
 - Time slice – each queue gets a certain amount of CPU time which it can schedule amongst its processes; *i.e.*, 80% to foreground in RR and 20% to background in FCFS.

Example:

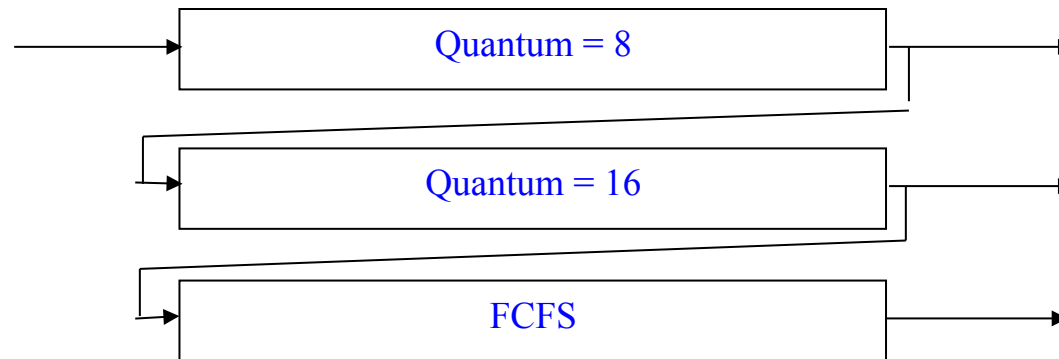


Multilevel Feedback Queue Scheduling

- ★ A process can move between the various queues
 - one form of aging mechanism to prevent starvation in multilevel queue scheduling.
- ★ Multi-level feedback queue scheduler is defined by the following parameters:
 - Number of queues.
 - Scheduling algorithm for each Queue.
 - Method used to determine when to upgrade a process to higher Queue.
 - Method used to determine when to demote a process to lower Queue.
 - Method for determining which Queue a process will enter when that process needs service.

Example

- ★ Three queues, Q_0 with time quantum = 8 ms, Q_1 with time quantum = 16 ms, and Q_2 with FCFS.
- ★ Scheduling
 - A new process enters Q_0 which is served FCFS. When it gains CPU, process receives 8 ms. If it does not finish in 8 ms, process is moved to Q_1 .
 - At Q_1 , process is again served FCFS and receives 16 additional ms. If it still does not complete, it is pre-empted and moved to Q_2 .



Guaranteed Scheduling

- ★ Make real promises to the user about performance, and then live up to them.
- ★ Example: for n users, the promise can be $1/n$ CPU time for each user.
- ★ System must keep track of how much CPU time a user has had for all his processes since login, and how long the user has logged in.

Multiple-Processor Scheduling

- ★ CPU scheduling is more complex when multiple CPUs are available.
- ★ *Homogeneous processors* within a multiprocessor → all processors are identical.
- ★ *Load sharing*:
 - A queue for each processor → load *not* balanced.
 - A single ready queue for all processors; two approaches:
 - ★ Each process is self-scheduling → need mutual exclusion on accessing the ready queue; symmetric multiprocessing.
 - ★ Master-slave structure → one processor as the scheduler; asymmetric multiprocessing.
- ★ *Asymmetric multiprocessing* – only one processor accesses the system data structures, alleviating the need for data sharing.
- ★ Asymmetric multiprocessing is simpler.

Real-time Scheduling

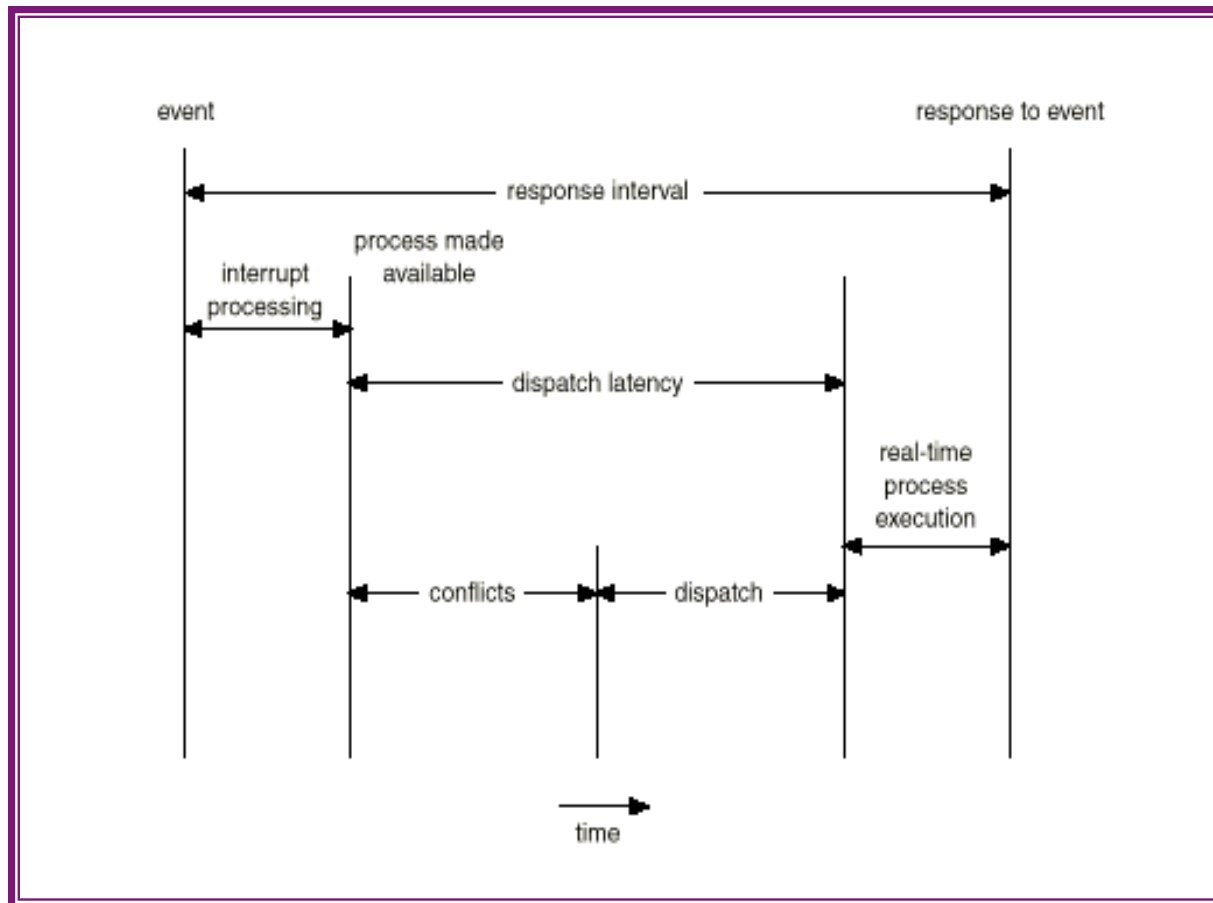
- ★ *Hard real-time* systems require to complete a critical task within a guaranteed amount of time.
 - Each submitted process has a statement of the amount of time needed to complete or perform I/O.
 - The scheduler:
 - ★ Admits the process if it can guarantee that the process will complete on time.
 - ★ Rejects the process if impossible.
 - The scheduler should know how long each type of OS function takes to perform → impossible in a system with secondary storage or virtual memory.
 - Hard real-time systems are composed of special purpose software running on hardware dedicated to their critical process.

Real-time Scheduling (cont.)

- ★ *Soft real-time* computing requires critical processes to have priority over others.
 - General-purpose system supporting multimedia, high- speed interactive graphics, *etc.*
 - Requirements:
 - ★ The system must have priority scheduling → real time processes always must have the highest priorities.
 - ★ Dispatch latency must be small; ways to achieve this goal:
 - Allow system calls to be preemptible.
 - Make the entire kernel preemptible → Create priority inversion: the high priority process waits for lower priority one to finish.
 - Solution: use priority inheritance protocol where a low priority task that is using resources needed by the higher priority task inherits its priority until completing the resources.

Real-time Scheduling (cont.)

- ★ Two components of conflict phase of dispatch latency:
 1. Preemption of a kernel process.
 2. Release of low-priority process resources that are needed by a high-priority process



Thread Scheduling

User-level thread scheduling

- * Thread library schedules user level threads to run on an available kernel level thread
 - Known as process-contention scope (PCS):
 - for many-to-one and many-to-many models
- * Priority scheduling is commonly used.
- * The thread is not actually running in CPU yet.
 - Kernel does not know the existence of a user-level thread.
 - Context switch (kernel) occurs when the time quantum for the process is up.

Kernel-level thread scheduling

- * kernel schedules which thread gets CPU
 - Known as system-contention scope (SCS)
 - For one-to-one model uses only SCS → Linux, Windows XP
 - The threads can be from the same or different processes.
- * Each thread is given time quantum, and context switch is for each thread.
- * Context switch in kernel thread is much more expensive than for user-level thread.
- * Context switch between threads from the same process is faster than from different processes.

Algorithm Evaluation

(a) **Deterministic model**

- Takes a particular predetermined workload and defines the performance of each algorithm for that workload.
- Simple, fast
- Requires exact number for input.
- Results are indicative only for this input.
- Too specific to be of general use.

Job	Burst time
1	10
2	29
3	3
4	7
5	12

Deterministic model Example

FCFS:

1	2	3	4	5
10	39	42	49	61

SJF(NP):

3	4	1	5	2
3	10	20	32	61

RR(Q=10):

1	2	3	4	5	2	5	2
10	20	23	30	40	50	52	61

Waiting Time Table:

Process	FCFS	SJF	RR
1	0	10	0
2	10	32	32
3	39	0	20
4	42	3	23
5	49	20	40
	140	65	115

Algorithm Evaluation (cont.)

(b) Queueing model

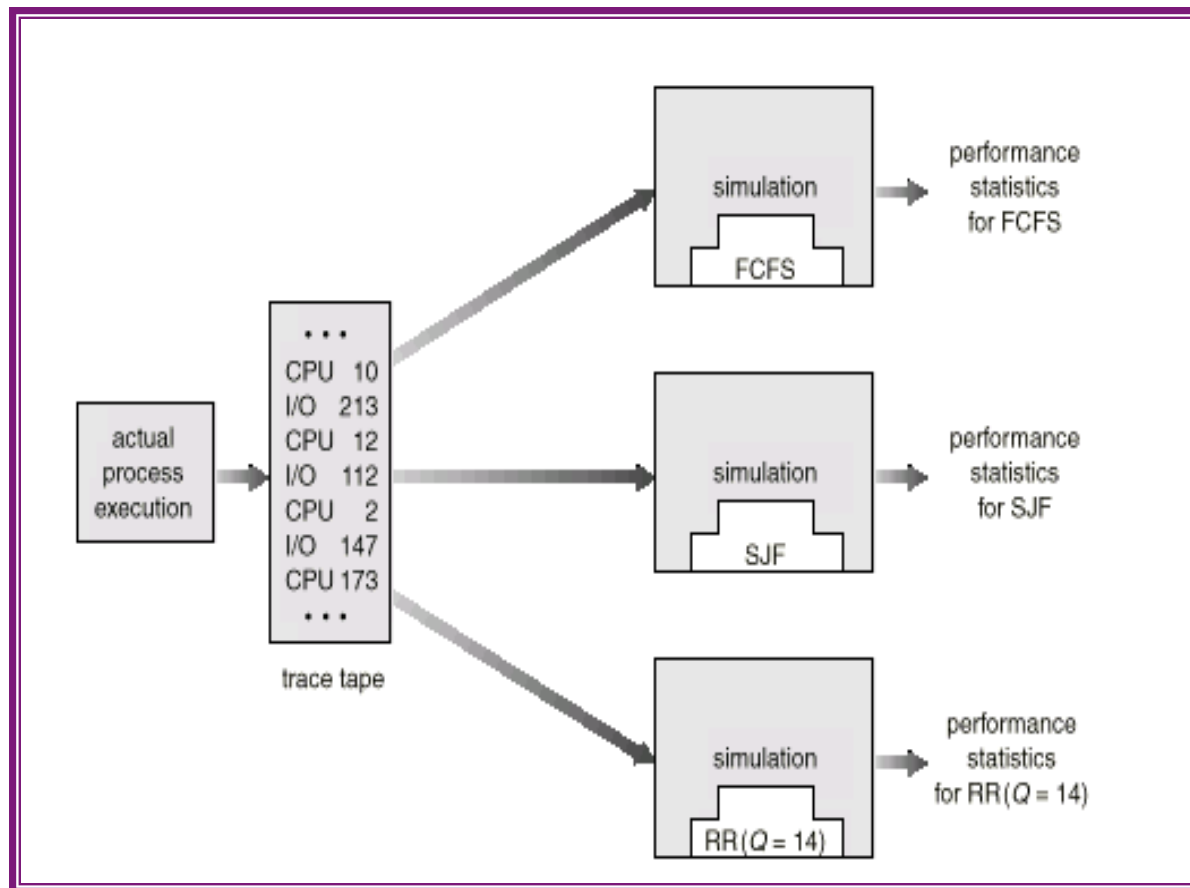
- Average queue length: n .
- Average waiting time in queue: w .
- Average arrival rate to the queue: λ .
- System steady state \rightarrow the number of processes leaving the queue must be equal to the number of processes that arrive, thus:
Little's formula: $n = \lambda * w$
- Queueing analysis is useful:
 - * In comparing algorithms.
 - * Only as good as the distributions (often difficult and unrealistic to make the problem mathematically tractable).

Algorithm Evaluation (cont.)

(c) Simulations.

- Involve programming a model of the computer system.
- Software data structures represent major components.
- Variable representing the clock.
- Data can be generated by:
 - * Random number generator.
 - * Trace tapes: created by monitoring a real system.
 - * Simulation can be very expensive:
 - Requiring hours of CPU time, large amount of storage (for trace tapes), and needs a lot time in designing, coding, and debugging the simulator.

Simulations (cont)



Algorithm Evaluation (cont.)

(d) Implementation.

- The only completely accurate way to evaluate a scheduling algorithm → code it, put it in the OS, see how it works.
- Difficulties:
 - ★ The cost of this approach.
 - ★ The environment in which the algorithm is used may change.

An Example: Linux 2.2

- ✱ Provides two separate process-scheduling algorithms:
 - Timesharing: fair preemptive
 - Real time: absolute priorities
- ✱ Linux 2.2 allows only user-mode processes to be pre-empted
- ✱ Each process has a scheduling class → a prioritised, credit-based algorithm
 - The first is for timesharing processes
- ✱ Prioritised, credit-based
 - Each process has a certain number of scheduling credits → the largest means the highest priority.
 - The running process' credit is decremented by one each time a timer interrupt occurs → the process is suspended when its credit becomes 0
 - If there is no runnable processes, do re-crediting to every process in the system:
$$\text{credits} = \text{credits}/2 + \text{priority}.$$
 - ✱ Requires $O(n)$ for this step, where n is the total number of processes.
 - Give high priority to interactive or I/O-bound processes.
- ✱ Real-time scheduling: FCFS, RR

Example $O(1)$ Scheduling

- ★ Scheduling problem with Linux 2.2:
 - Not adequately support SMP systems.
 - Getting slower when the total number of tasks n increases; $O(n)$
- ★ Scheduling algorithm in Linux 2.5 is $O(1)$, and provides better support for SMP systems.
 - The algorithm runs in constant time, irrespective of the total number of processes n
- ★ Preemptive and priority based
- ★ Two separate priority range (lower value has higher priority):
 - For real time: 0 to 99
 - Nice: 100 to 140.
- ★ Higher priority task is assigned with higher time quantum, *i.e.*,
 - A process with priority 0 $\rightarrow q = 200\text{ms}$.
 - A process with priority 140 $\rightarrow q = 10\text{ms}$.

$O(1)$ Scheduling (cont.)

- * Kernel maintains a list of runnable tasks in a runqueue that contains two priority arrays
 - Active:** contains all tasks with time remaining in their time slices.
 - Expired:** contains all tasks with no remaining time in their time slices.
- * Scheduler selects the highest priority task in Active array, indexed by priority.
- * For SMP, each processor maintains its own runqueue and schedules itself independently
- * Recalculate new priority of each process with exhausted time and move it to **Expired**
- * Calculate its new priority:
 - Real time task is set with fix priority.
 - All other tasks have dynamic priorities (calculated before going to **Expired**):
 $nice \text{ value} + d$; where $-5 \leq d \leq +5$
 - The value for d is determined by the task interactivity: How long the process has been sleeping waiting for I/O.
 - * Longer sleep time \rightarrow more interactive $\rightarrow d$ is set close to -5.
 - CPU bound process \rightarrow gets lower priority.
- * When Active queue is empty \rightarrow swap Active and Expired.

Completely Fair Scheduler (CFS)

- ✱ $O(1)$ has poor response time for interactive processes
- ✱ Kernel release 2.6.23 makes CFS as the default Linux scheduler
- ✱ Scheduling in Linux is based on scheduling classes
 - Each class has a specific priority
 - The next task to run is the task with the highest priority in the highest priority class.
- ✱ Standard Linux implements two scheduling classes (A new scheduling class can be added)
 - A default class using CFS scheduler
 - A real time scheduling class
- ✱ CFS assigns a proportion of CPU time to each task
 - The proportion is calculated based on the *nice* value (-20 to +19) assigned to each task
 - ✱ Lower value means higher priority → receives higher proportion of CPU time
 - ✱ The default *nice* value is 0

CFS (cont.)

- ★ CFS uses *targeted latency*: an interval of time during which every runnable task should run at least once.
 - Targeted latency has default and minimum values
 - Targeted latency increases when the total number of tasks in the system increases above a threshold value
 - Proportions of CPU time are allocated from the value of targeted latency
- ★ CFS keeps virtual run time (vruntime – how long a task has run) for each task
- ★ Vruntime is associated with decay factor based on priority of the task
 - Lower priority has higher decay factor
 - Normal task (*nice* value 0) has vruntime equal to physical runtime
 - ★ A normal task that runs for 200 milliseconds has vruntime = 200 milliseconds
 - ★ A lower priority task that runs for 200 milliseconds has vruntime > 200 milliseconds
 - ★ A higher priority task that runs for 200 milliseconds has vruntime < 200 milliseconds
 - CFS selects the task with the smallest vruntime to run next
 - ★ A higher priority task can preempt a lower priority task