

COMMONWEALTH OF AUSTRALIA
Copyright Regulation 1969

WARNING

This material has been copied and communicated to you by or on behalf
of **Curtin University of Technology** pursuant to Part VB of the
Copyright Act 1968 (**the Act**)

The material in this communication may be subject to copyright under the
Act. Any further copying or communication of this material by you
may be the subject of copyright protection under the Act.

Do not remove this notice

Operating Systems

COMP2006

Processes and Threads

Lecture 2

Process and Thread

References: Silberschatz, Galvin, and Gagne, *Operating System Concepts*, Chapters 3, 4

Topics:

- ★ Process concept.
- ★ Operations on processes
- ★ Inter-process communication.
- ★ Threads.

Process Concept

- ★ Multiprogramming allows multiple programs to be loaded into memory and to be executed concurrently
 - Need stricter control on the various programs.
- ★ In a batch system, we call the program as *jobs*;
 - In a time-shared system, we call them *user programs* or *tasks*;
 - Textbook uses the terms *job* and *process*.
- ★ *Process* is a program in execution;
 - A program (executable file containing list of instructions) is passive
 - ★ It does nothing unless its instructions are executed.
 - A process has a Program Counter (PC)
 - ★ It specifies the next instruction to execute
 - ★ Process execution progresses sequentially – the CPU executes one instruction of the process after another until the process completes.
 - A process needs resources to accomplish its task
 - ★ Resources, e.g. CPU time, memory, files, and I/O devices.
 - ★ Several processes may execute concurrently by multiplexing CPU among them.

Process Concept (cont.)

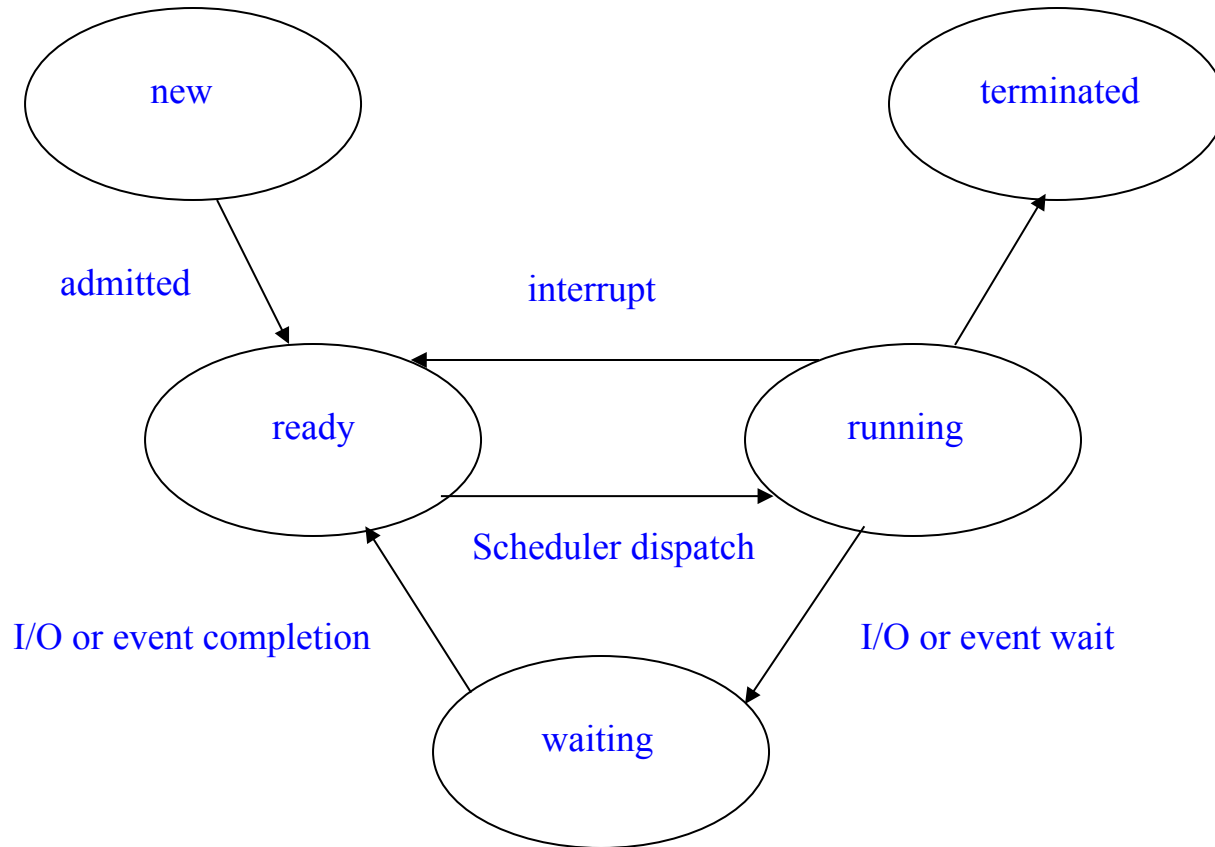
- ★ A process includes:
 - Text section: the program code
 - Program counter: contains a pointer to the next instruction to execute
 - Contents of processor registers.
 - Stack: contains temporary data; e.g., subroutine parameters, return addresses, temporary variables.
 - Data section: contains global variables.
 - Heap: memory that is dynamically allocated during run time
- ★ Two processes may be associated with the same program
 - They are considered two separate execution sequences.
- ★ Process Management of an OS is responsible for the following:
 - Creation and deletion of user and system processes.
 - Suspension and resumption of processes.
 - Provision of mechanisms for process synchronisation.
 - Provision of mechanisms for process communications.
 - Provision of mechanisms for deadlock handling.

Process Concept (cont.)

- ★ A process can be **user** level or **kernel** level
- ★ As a process executes, it changes state.
- ★ Each process may be in one of the following process states:
 - ★ *New* – the process is being created.
 - ★ *Running* – instructions are being executed.
 - ★ *Waiting* – the process is waiting for some event to occur.
 - ★ *Ready* – the process is waiting to be assigned to a processor.
 - ★ *Terminated* – the process has finished execution

Process Concept (cont.)

Transition diagram of process state:



Process Concept (cont.)

- ★ Each process is represented by a data structure
 - called Process Control Block (PCB) or Task Control Block
- ★ Each PCB contains:
 - Process number (pid).
 - Process state: new, ready, running, waiting, etc.
 - Program counter: shows the address of next instruction to be executed.
 - CPU registers: the registers vary in number and type, depending on the computer architecture.
 - CPU scheduling information: process priority, pointers to schedule queues, etc.
 - Memory management information: base and limit register, page tables, etc.
 - Accounting information: amount of CPU and time used, etc.
 - I/O status information: the list of I/O devices allocated to this process, a list of open files, etc.

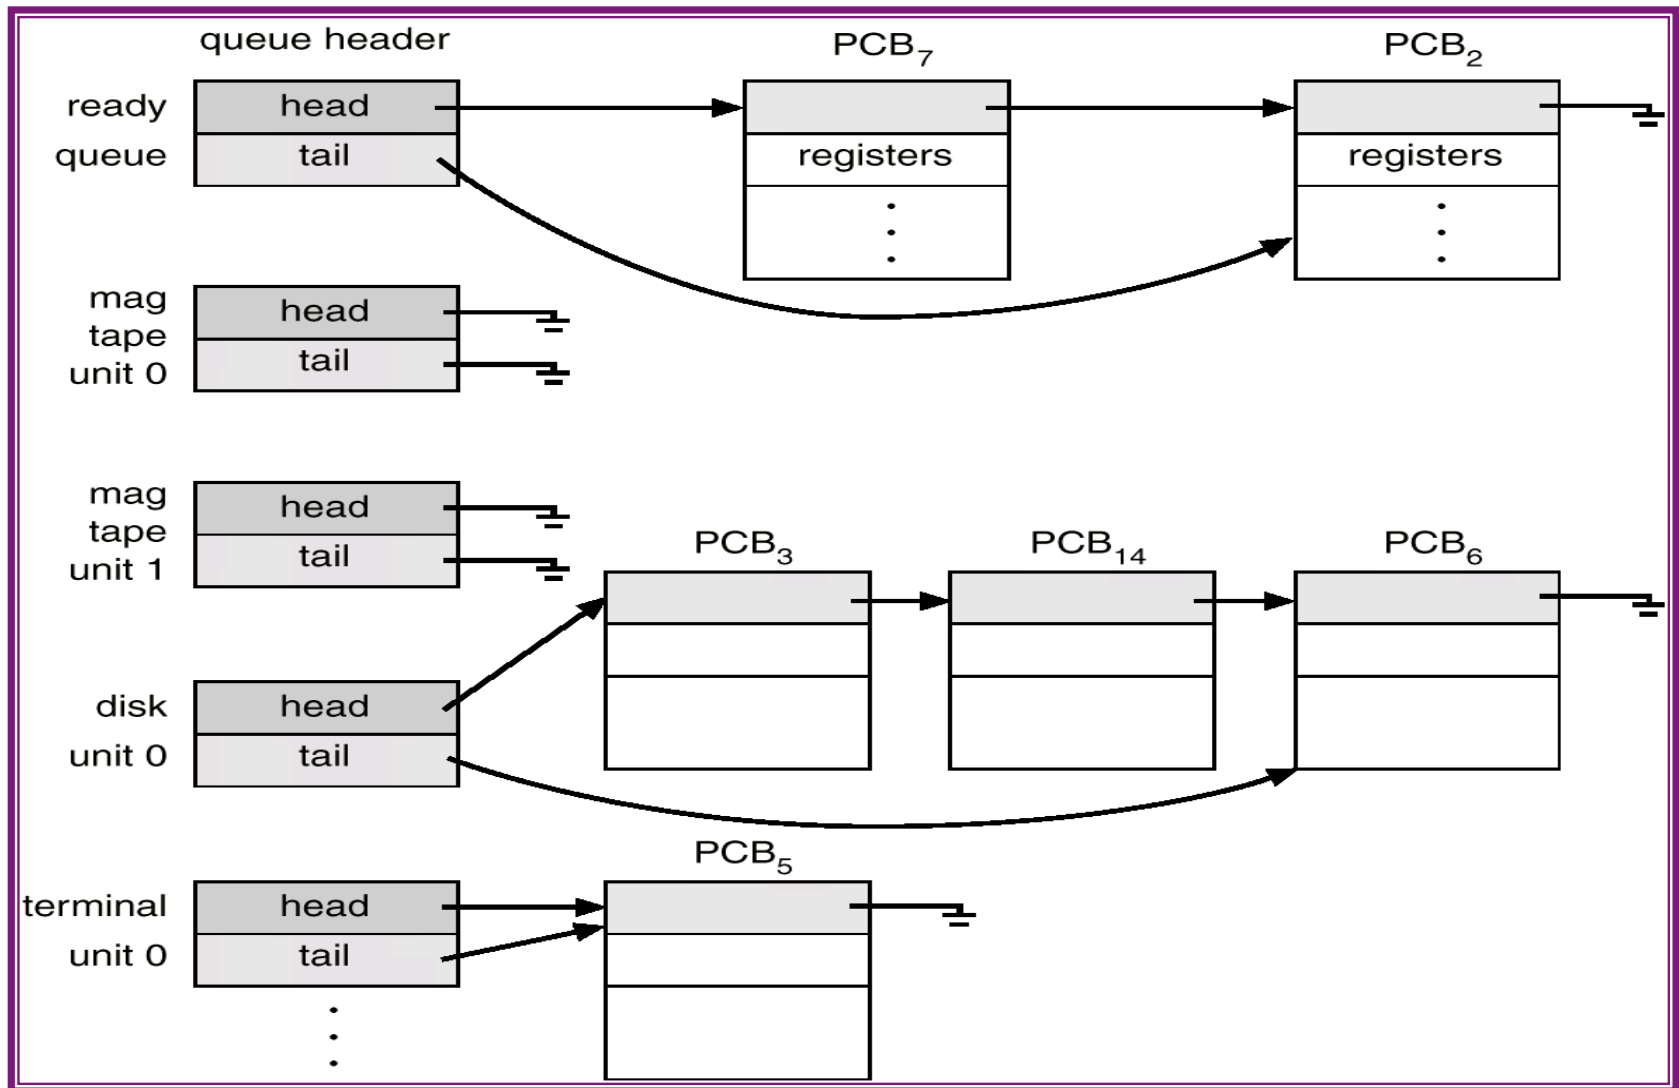
Process Concept (cont.)

- ★ PCB in Linux is a C structure *task_struct* that contains:
 - long state; // *state of the process*
 - struct sched_entity se; // *scheduling information*
 - struct task_struct *parent; // *the parent of this process*
 - struct task_struct *children; // *the children of this process*
 - struct file_struct *files; // *list of open files*
 - struct mm_struct *mm // *this process's address space*
- ★ **Parent** of a process is a process that creates it
 - Its **children** are processes that it creates
 - **Siblings** are children with the same parent
- ★ All active processes are represented by a double linked list of *task_struct*
 - A variable *current* is used as a pointer to currently executing process.

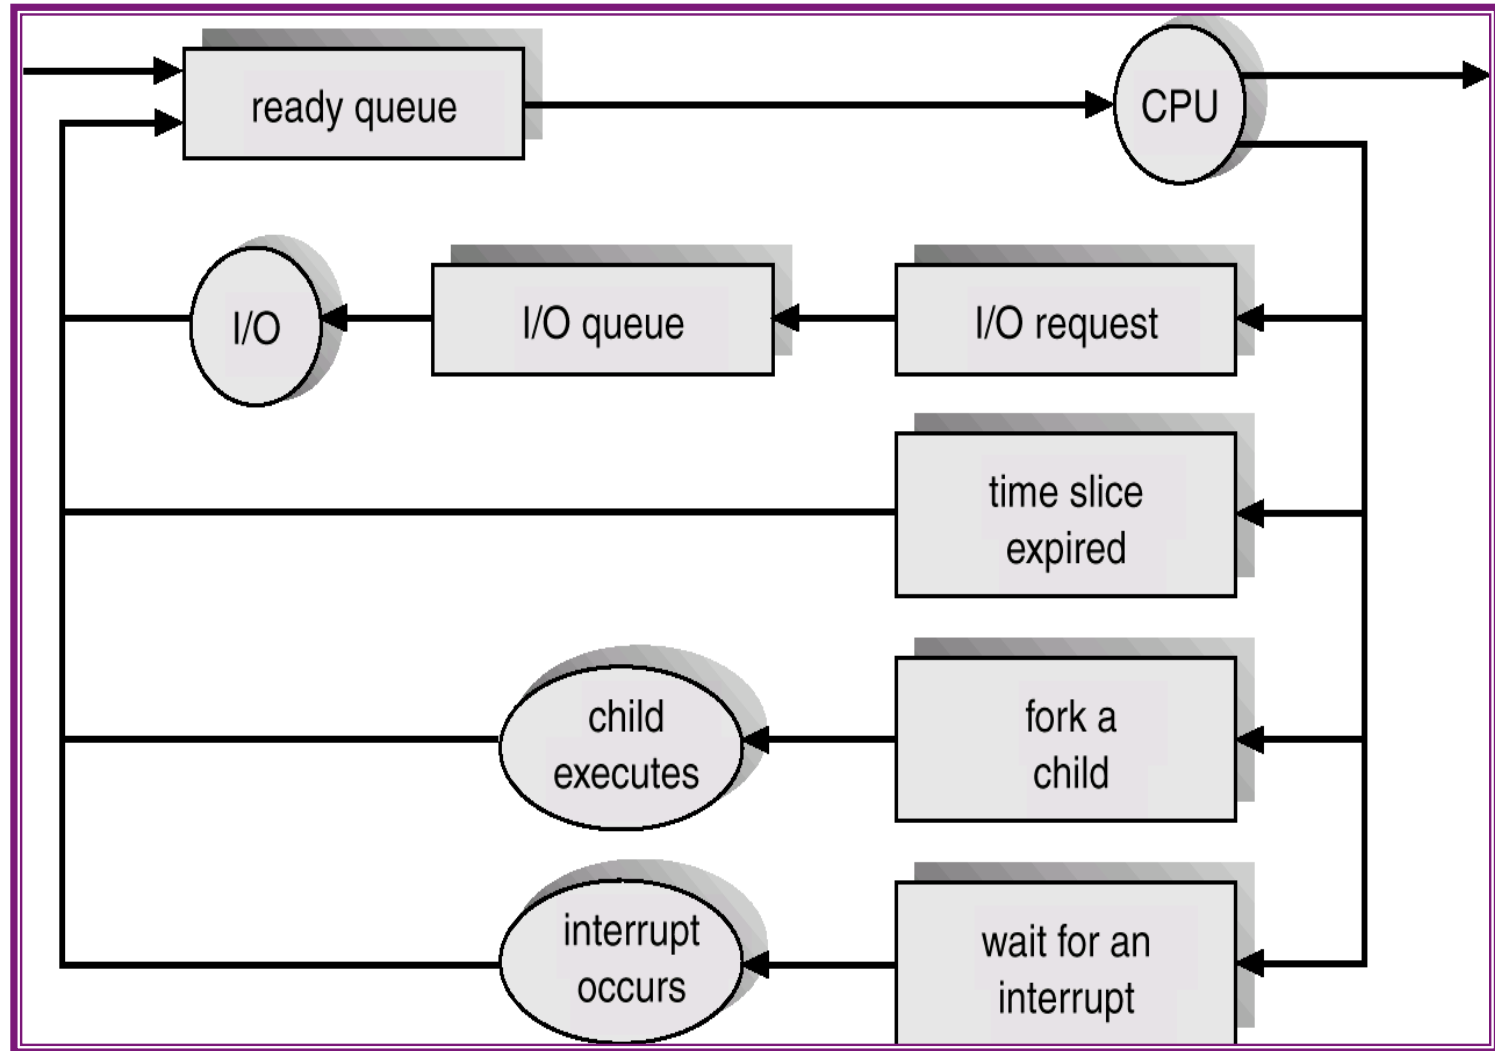
Process Scheduling

- ★ OS uses process/CPU scheduler to select one of available processes to be executed by CPU
 - The scheduler is used to meet system objective
 - ★ Multiprogramming aims to maximize CPU utilization
 - ★ Timesharing aims to switch processes frequently so that users can interact with their running programs
- ★ OS keeps several queues (implemented as link lists):
 - **Job queue** – All processes entering the system are put in a job queue.
 - **Ready queue** – A set of all processes residing in main memory, ready and waiting to execute.
 - **Device queues** – A set of processes waiting for an I/O device
 - ★ Each device has its own device queue.
- ★ Process migrates between the various queues
 - Each queue is a list of PCBs
 - When a process terminates, the OS removes its PCB from the queue and de-allocates its resources.

Ready Queue and Various I/O Device Queues

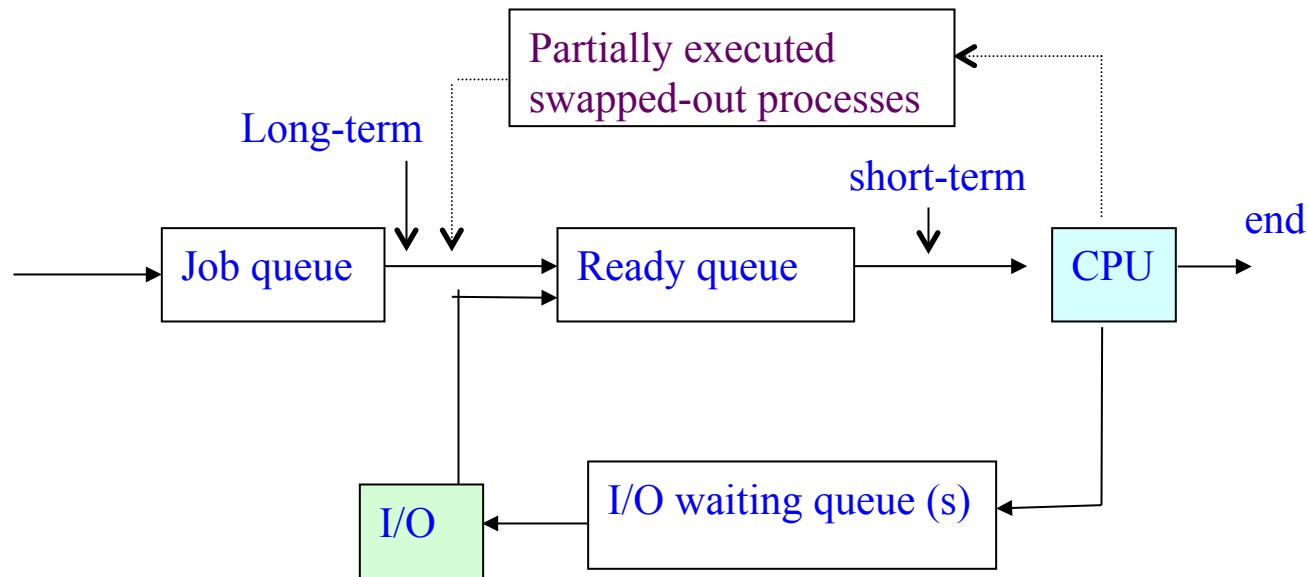


Representation of Process Scheduling



Schedulers

- ★ A process migrates between various scheduling queues.
 - Selection on which process to migrate is done by the appropriate scheduler.
- ★ Long-term scheduler (or job scheduler) selects which processes should be brought into the ready queue.
 - It is invoked very infrequently (seconds, minutes) → the scheduler can be slow
 - It controls the degree of multiprogramming (the number of processes in memory).
 - A time sharing system often has no long-term scheduler.



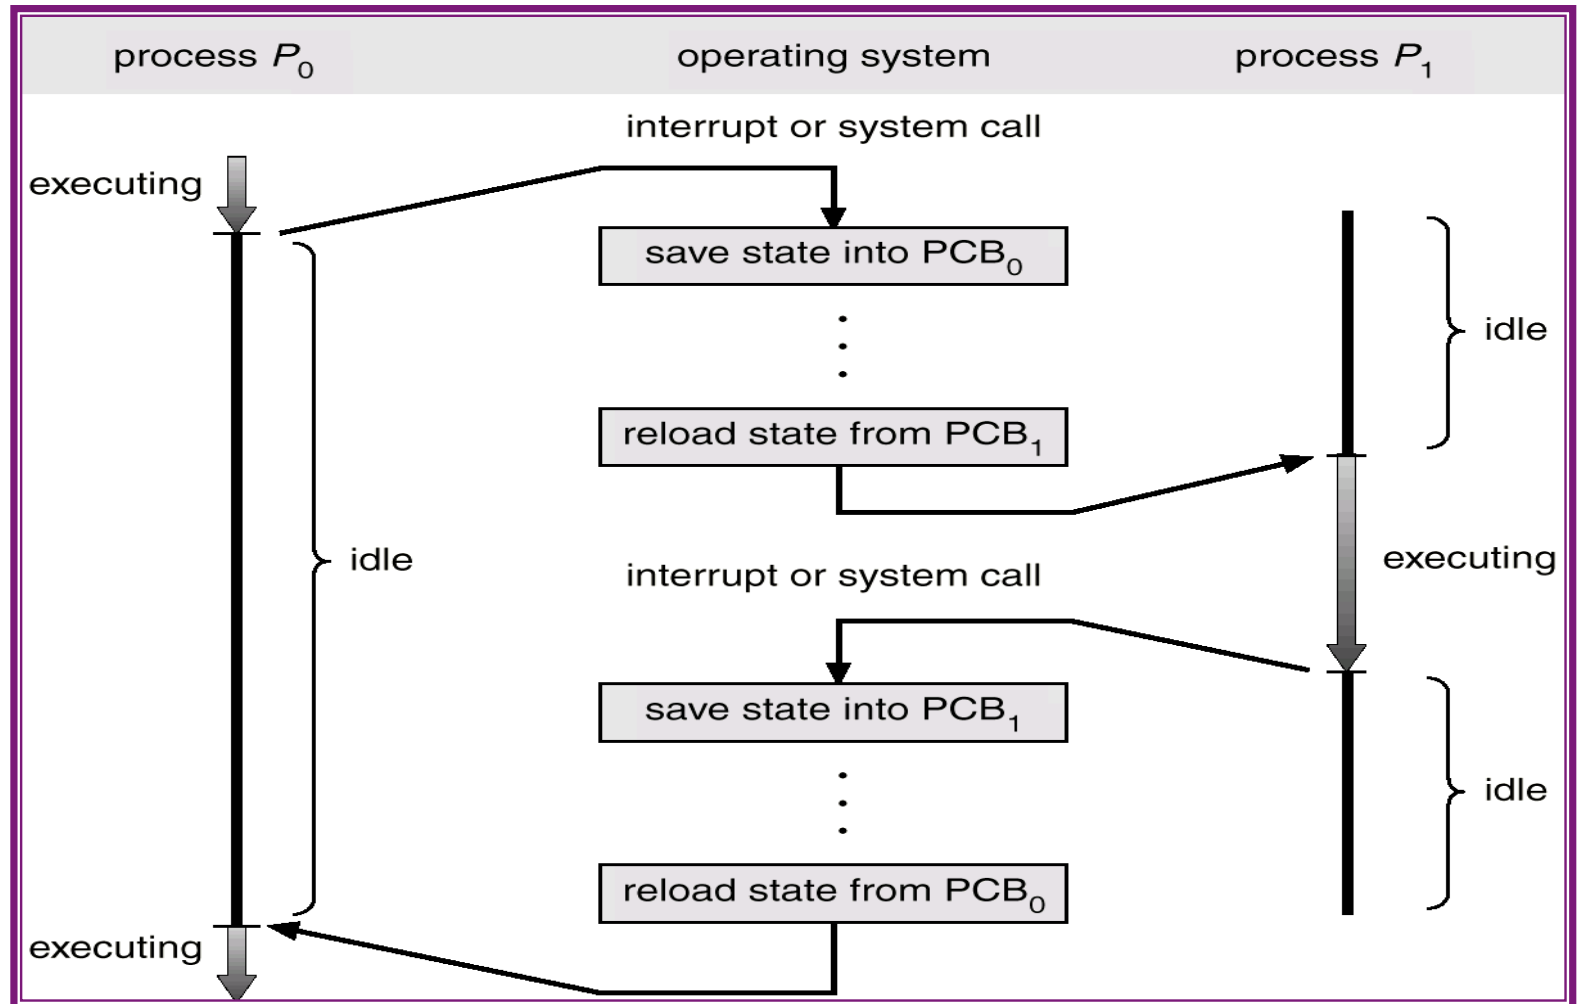
Schedulers (cont.)

- ★ Short-term scheduler (or CPU scheduler) selects or **dispatches** a process to be executed next and allocated CPU
 - It is invoked very frequently (milliseconds); thus the scheduler must be fast.
- ★ A process can be described as either:
 - I/O-bound process
 - ★ It spends more time doing I/O than computations
 - ★ It has many short CPU bursts.
 - CPU-bound process
 - ★ It spends more time doing computations
 - ★ It has few very long CPU bursts.
- ★ Long-term scheduler needs to select a good mix of I/O bound and CPU bound processes.
- ★ Medium-term scheduler is sometimes introduced to remove processes from memory to reduce the degree of multiprogramming.

Context Switch

- ★ When CPU switches to another process, the system must save the *context* of the old process and load the *context* for the new process
 - The OS suspends the old process and runs the new one
 - The saved context is needed when the old process resume its execution
- ★ Context-switch time is overhead
 - The system does no useful work while switching
 - Performance bottleneck.
- ★ Context-switch-time is dependent on hardware support
 - Typically from a few milliseconds.
- ★ When do we switch CPU?
 - Job voluntarily waits (system calls).
 - Interrupt: Higher priority event/job needs attention.
 - Interrupt: Timer.

CPU Switch From Process to Process



Operations on processes

Process Creation

- * Parent process creates child processes, which in turn can create other processes, forming a tree of processes
 - in Unix use system call `fork()`.
 - Each process is known by its process ID (an integer value)
 - * The init process has `PID = 1`; it is the root parent process for all user processes
 - * Use `ps -el` command to list all processes active in the system
 - A process needs resources (CPU, memory, files, I/O, etc.).
- * When a process creates a sub-process:

For resources


- Parent and child processes may share all resources.
- Child processes may share subset of parent's resources.
- Parent and child processes may share no resources.

For execution

- Parent and child processes may execute concurrently → Linux
- Parent may wait until child process terminates
 - * parent calls `wait()` to move itself out from ready queue


For address space

- Child process address space is duplicate of parent's.
- Child has a program loaded into it → in UNIX: use `exec()`.

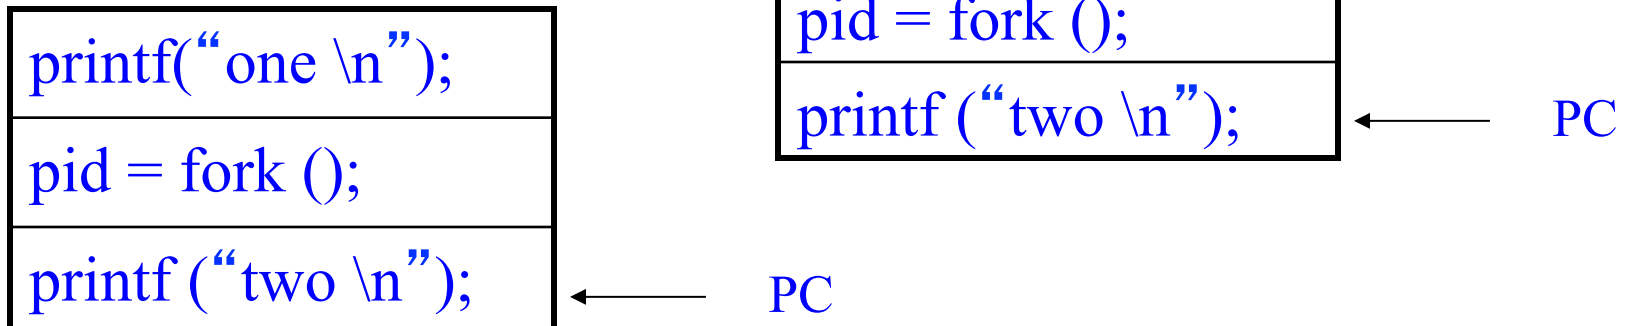


UNIX example

- ★ **fork ()** system call creates a new process.



After executing fork ():



Operations (cont.)

```
int main ()
{
    pid_t pid;

    pid = fork (); // create a child process

    if (pid < 0) { // error occurs
        return 1;
    }
    else if (pid == 0) { // child process; OS returns 0 to the child process
        execute_function_for_child(); // alternatively, use execlp ("/bin/ls", "ls",
                                     NULL) to execute the command "ls"
    }
    else { // pid > 0 is the child's PID returned to the parent process
        wait (NULL); // the parent process waits for its child's termination
    }
}
```

Operations (cont.)

Process termination

- ★ A process terminates when it executes last statement and/or calls **exit ()**;
 - OS will deallocate the child's resources
 - the child can return its status value to parent that calls **pid = wait (&status)**
 - ★ If the parent does not call wait(), the OS does not remove the child from process table
 - The child process becomes a '*zombie*'
 - When the parent terminates, the child becomes *orphan*, and *init* will automatically be its parent.
 - *init* process periodically calls wait() removing the orphans from process table
- ★ Parent may terminate execution of child processes (**abort**); Reasons for termination:
 - Child has exceeded allocated resources.
 - Task assigned to child is no longer required.
 - Parent is exiting
 - ★ Some OS does not allow child to continue if its parent terminates.
 - ★ Linux does not terminate the child when its parent terminates

Cooperating Processes

- ★ Concurrent processes executing in OS may be either independent or cooperating processes.
 - *Independent processes* cannot affect or be affected by the execution of another process.
 - *Cooperating processes* can affect or be affected by the execution of another process.
- ★ Advantages of process cooperation:
 - Information sharing
 - ★ Several users may need the same piece of information.
 - Computation speed-up
 - ★ Break a task into subtasks and execute in parallel.
 - Modularity
 - ★ Divide the system function into separate processes.
 - Convenience
 - ★ A user may want to do editing, printing in parallel.
- ★ Cooperating processes can communicate via:
 - **shared-memory** (Read Section 3.5.1 POSIX Shared memory)
 - **message passing**

Shared memory vs. message passing

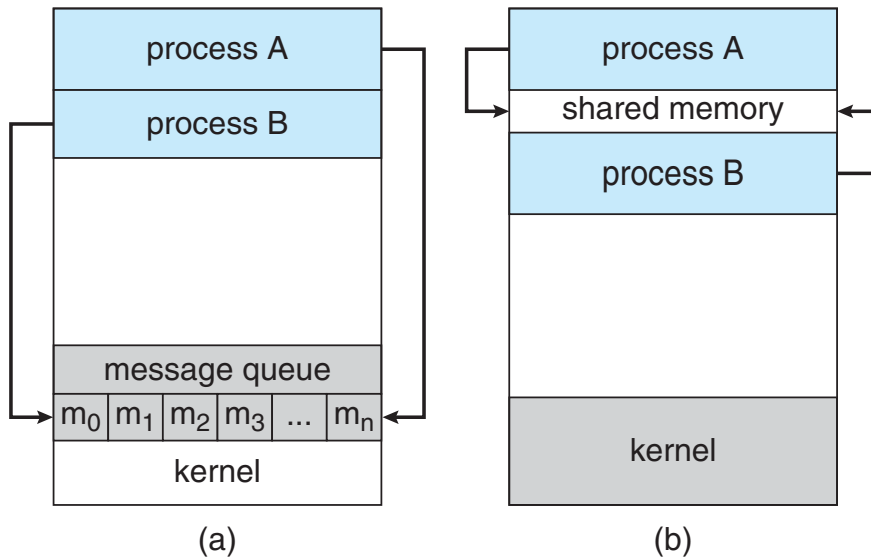


Fig. 3.12 (textbook)

Shared memory

- * A process uses a system call to create a shared memory region
 - Other process that wants to communicate via the memory region must attach to it to its address space (use a system call)
 - Once established, all accesses to the shared memory are as if accessing normal memory area (no system call)
 - However, all accesses must be synchronized

Message Passing

- * Use a system call to exchange message between processes
 - More time consuming
 - OS handles synchronization between processes
 - Good for distributed system with small amount of data exchanged
 - Good for multicore system
 - * Shared memory suffers from cache coherency issues.

Example of concurrent processes

Producer-consumer problem

- ★ Paradigm for cooperating processes: *producer* process produces information that is consumed by a *consumer* process.
- ★ To allow producer and consumer processes to run concurrently, there is a buffer of items that can be filled by the producer, and emptied by the consumer.
 - *Unbounded-buffer* places no practical limit on the size of the buffer.
 - *Bounded-buffer* assumes that there is a fixed buffer size.
- ★ The buffer can be in a form of shared memory, or provided by the OS via message passing.

Producer-consumer (cont.)

Bounded-buffer with shared memory solution

/* Shared data */

var *buffer*: **array** [0..*n*-1] of *item*;

in, *out*: 0..*n*-1; **/* initially $in = out = 0$; in is the pointer to the next free position; out is the pointer to the first full position */**

Producer process

repeat

 produce an item in *nextp*;

 ...

// while buffer is full

while ($in+1 \bmod n == out$) **do no-op**;

buffer[*in*] = *nextp*;

in = $in+1 \bmod n$;

until false;

Consumer process

repeat

// while buffer is empty

while ($in == out$) **do no-op**;

nextc = *buffer*[*out*];

out = $out+1 \bmod n$;

 ...

 consume the item in *nextc*;

 ...

until false;

- Solution is correct, but can only fill up $n-1$ buffer;
 - How to fill up n items in the buffer?

Message Passing

- ★ Message passing facility provides at least two operations:
 - **Send** (message) – message size fixed or variable.
 - **Receive** (message).
- ★ If P and Q wish to communicate, they need to:
 - Establish a *communication link* between them.
 - Exchange messages via *send/receive*.
- ★ Several methods for logically implementing a link and *send/receive* operations:
 - Direct or indirect communication.
 - Synchronous or asynchronous communication.
 - Automatic or explicit buffering.

Message Passing (cont.)

Direct communication

- ★ In the direct-communication, each process must explicitly name the recipient or sender of the communication.
- ★ Primitives used:
 - **Send** (P, message) – send a message to process P.
 - **Receive** (Q, message) – receive a message from process Q.
- ★ Properties of communication link for this scheme:
 - Links are established automatically.
 - A link is associated with exactly one pair of communicating processes.
 - Between each pair there exists exactly one link.
 - The link may be uni-directional, but is usually bi-directional.

Example

Producer process

repeat

...

produce an item in *nextp*;

...

send (*consumer*, *nextp*)

until *false*;

Consumer process

repeat

receive (*producer*, *nextc*);

...

consume the item in *nextc*;

...

until *false*;

Message passing (cont.)

Indirect communication

- ★ Messages are sent to and received from mailboxes.
 - Each mailbox has a unique id.
 - Processes can communicate only if they share a mailbox.
- ★ Primitives used:
 - **Send** (A, message): send a message to mailbox A.
 - **Receive** (A, message): receive a message from mailbox A.
- ★ Properties of communication link for this scheme:
 - Link established only if processes share a common mailbox.
 - A link may be associated with many processes.
 - Each pair of processes may share several communication links.
 - Link may be uni-directional or bi-directional.
- ★ Operations:
 - Create a new mailbox.
 - Send and receive messages through mailbox.
 - Destroy a mailbox.

Indirect communication(cont.)

Mailbox sharing

- ★ Consider processes P1, P2, and P3 share a mailbox A, and process P1 sends a message to mailbox A while processes P2 and P3 execute a **receive** from A,
 - Which process will receive the message sent by P1?
- ★ Three possible solutions:
 - Allow a link to be associated with at most two processes.
 - Allow only one process at a time to execute a receive operation.
 - Allow the system to select the receiver arbitrarily. Sender is notified who the receiver was.

Synchronization

- ★ Message passing can be either *blocking* (*synchronous*) or *nonblocking* (*asynchronous*).
- ★ Blocking send: The sender is blocked until the message is received.
- ★ Nonblocking send: the sender resumes operation after sending the message.
- ★ Blocking receive: the receiver blocks until message is available.
- ★ Nonblocking receive: receiver retrieves either a valid message or a NULL.

Buffering

- ★ A link has some capacity that determines the number of messages that can reside in it temporarily
- ★ This can be viewed as queue of messages attached to the link; implemented in one of three ways:
 - zero capacity – 0 messages: sender must wait for the receiver (rendezvous); called as no-buffering.
 - bounded capacity – finite length of n messages: sender must wait if link is full.
 - unbounded capacity – infinite length: sender never waits.
- ★ For the non-zero capacity (automatic buffering), the sender does not know if a message has arrived at the destination.
 - may use asynchronous communication; the receiver sends back an acknowledgement to the sender.

Exception conditions

- ✱ Error may happen during process communication;
- ✱ When a failure occurs either in centralized (single machine) or distributed (processes reside in different machine), some error recovery must take place.

- Process terminates; e.g., process P is waiting for a message from a terminated process Q.

Solution: OS terminates P or notify P that Q has terminated.

- Lost messages; e.g., message from P to Q become lost due to hardware error;
 - ✱ The most common lost detection method is by using timeout
 - ✱ An acknowledgement must be received by sender.

Solution:

- ✱ OS detects it and resends the message.
 - ✱ Sender detects it and resends.
 - ✱ OS detects it and lets the sender know.
- Scrambled messages – message is received but in error
 - ✱ It is detected by parity checking, CRC.


Solution: resend it.

Threads

- ★ A *thread* of control is an independent sequence of execution of program code in a process.
- ★ A thread (or lightweight process) is a basic unit of CPU utilization.
 - A traditional process (or heavyweight process) is equal to a task with one thread.
- ★ A traditional process has a single thread that has sole possession of the process's memory and other resources
 - Context switch becomes performance bottleneck
 - Threads are used to avoid the bottleneck.
 - Threads share all the process memory, and other resources.
- ★ Threads within a process:
 - are generally invisible from outside the process.
 - are scheduled and executed independently in the same way as different single-threaded processes.
- ★ On a multiprocessor, different threads may execute on different processors
 - On a uni-processor, threads may interleave their execution arbitrarily.

Threads (cont.)


Three processes with one thread each



Process

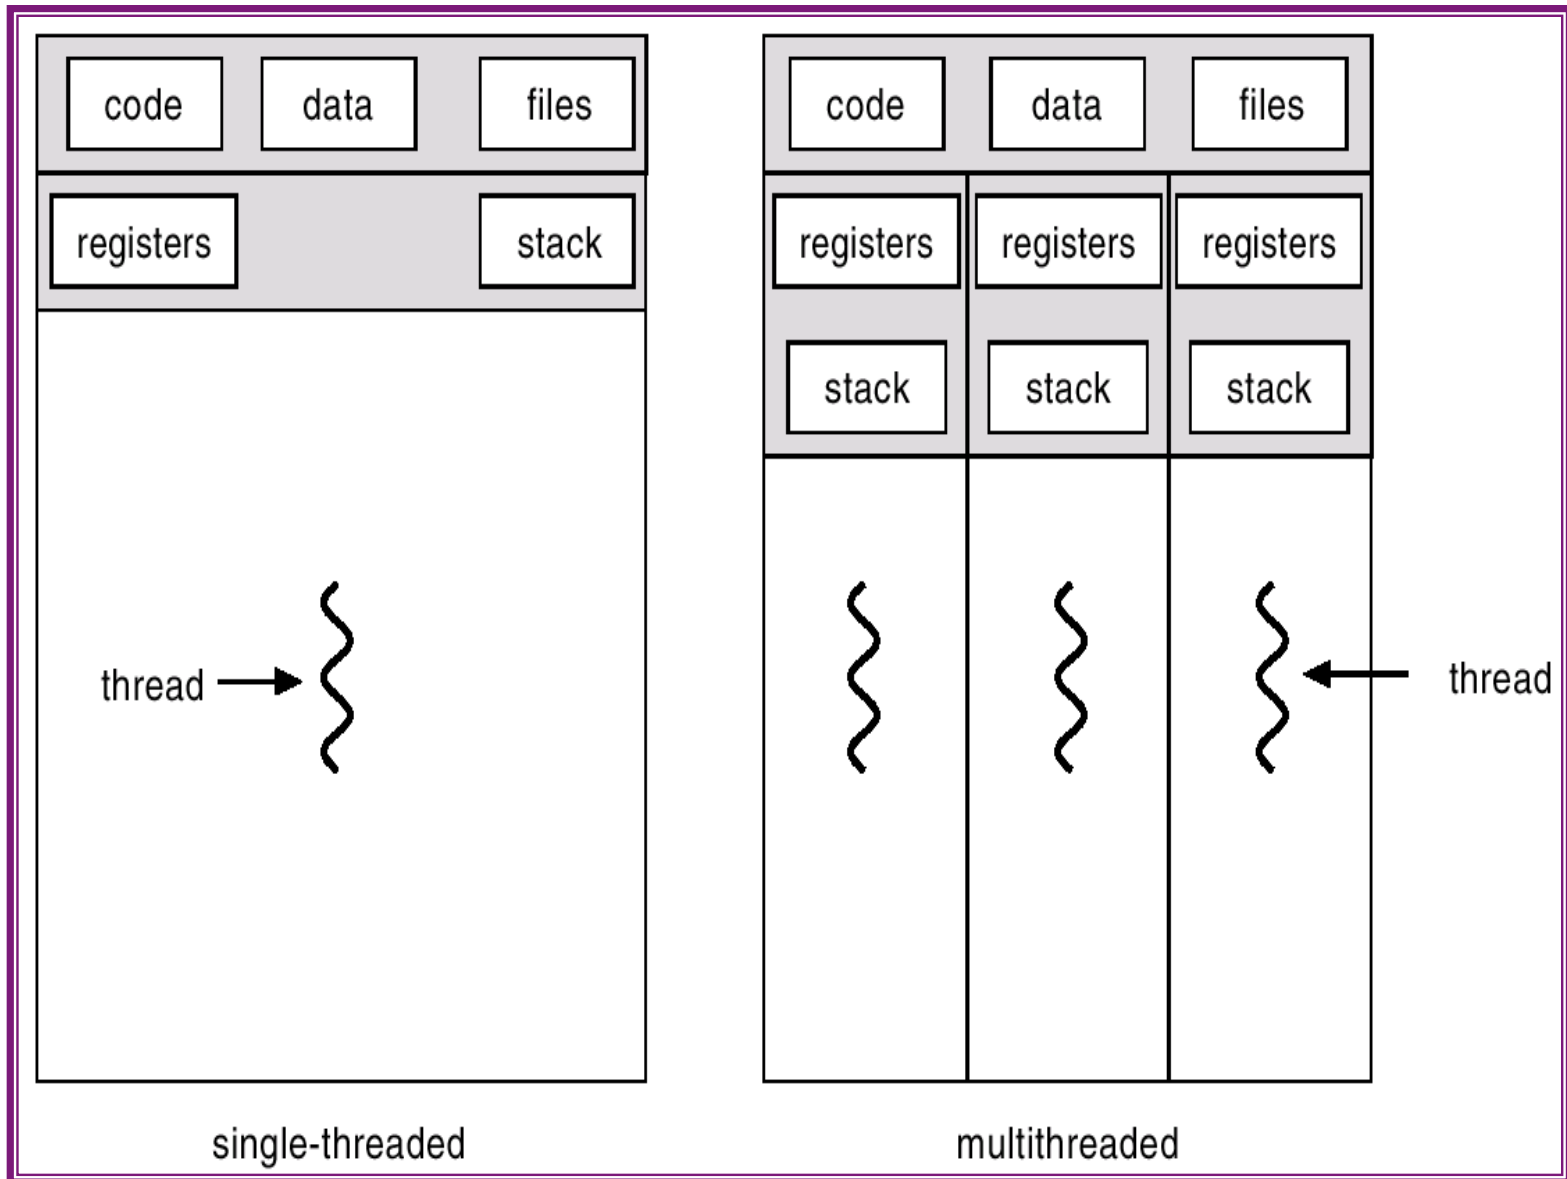
Thread

PC



One process with three threads

Single and Multithreaded Processes



Threads (cont.)

- ★ Threads operate, in many respects, in the same manner as processes:
 - Threads can be in one of several states: ready, blocked, running, or terminated, etc.
 - Threads share CPU
 - ★ Only one thread at a time is running.
 - A thread within a process executes sequentially, and each thread has its own PC and Stack.
 - Thread can create child threads, can block waiting for system calls to complete
 - ★ If one thread is blocked, another can run.
- ★ One major different with process: threads are not independent of one another
 - All threads can access every address in the task
 - ★ A thread can read or write any other thread's stack.
 - There is no protection between threads (within a process)
 - ★ However this should not be necessary since processes may originate from different users and may hostile to one another while threads (within a process) should be designed (by same programmer) to assist one another.

Threads (cont.)

Benefits

- ★ **Responsiveness:** For interactive application, when one thread in the program is blocked and waiting, a second thread in the same task can run.
- ★ **Resource Sharing:** By default, threads share the memory and resources within a process
 - Applications that require sharing a common buffer (i.e., producer-consumer) benefit from thread utilization.
- ★ **Economy:** It is more economical to create and context switch threads than processes;
 - in Solaris 2, creating a process is about 30 times slower than creating a thread, and context switch is about five times slower.
- ★ **Utilization of multiprocessor architectures:** Each thread may run in parallel on a different processor.

Threads (cont.)

User and Kernel Threads

- ★ Threads can be:
 - Kernel-supported threads.
 - User-level threads.
- ★ In kernel-supported threads, a set of system calls similar to those for processes are provided.
 - Supported by the Kernel;
 - ★ Kernel does context switch from one thread to another → time consuming.
 - Windows XP, Mac OS X, Solaris 2, Tru64 UNIX, BeOS, Linux.

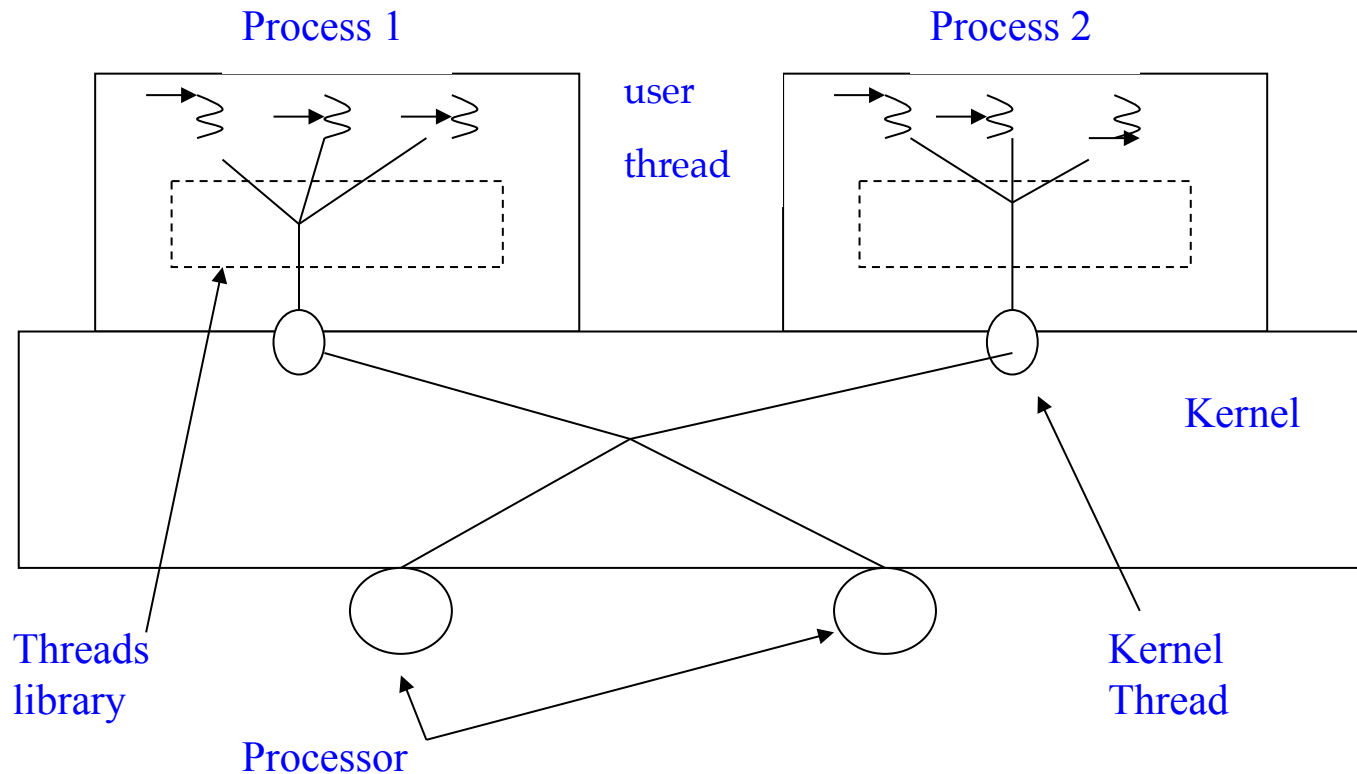
Threads (cont.)

- ★ User-level threads: supported above the kernel, via a set of library calls at the user level.
 - Kernel does not know that there are user-level threads.
 - Thread management is done by user-level threads library.
 - Faster than kernel-level.
 - Drawback: Kernel considers a set of user-level threads as a single thread;
 - ★ if any user level thread in the set is blocked (e.g., I/O), then other threads in the set cannot run.
 - POSIX Pthreads, Mach C-threads, Solaris 2 UI-threads.
- ★ Hybrid approach implements both user-level and kernel-supported threads (Solaris 2).

User to Kernel Threads Mapping

Many-to-one Model:

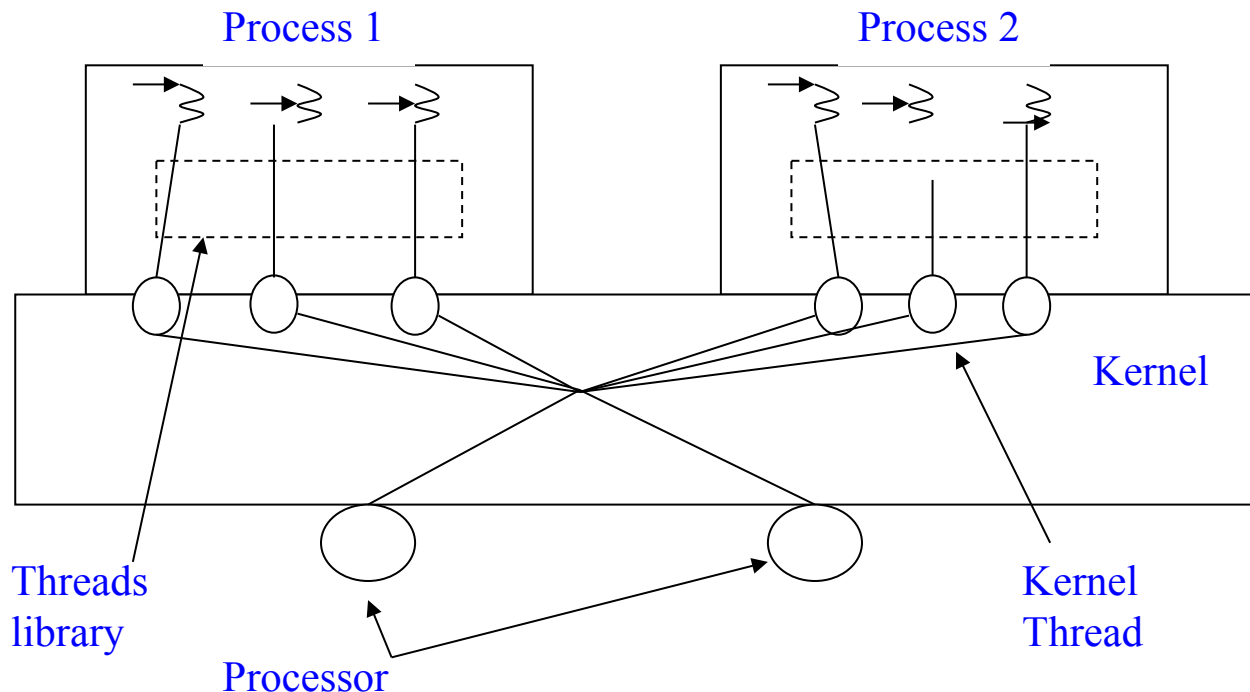
- ★ Many user-level threads are mapped to one kernel thread.
- ★ Multiple threads within a process cannot run in parallel on multiprocessors.
- ★ Used by Solaris 2 (Green threads), and threads libraries for systems with no kernel-threads.



Mapping (cont.)

One-to-one Model:

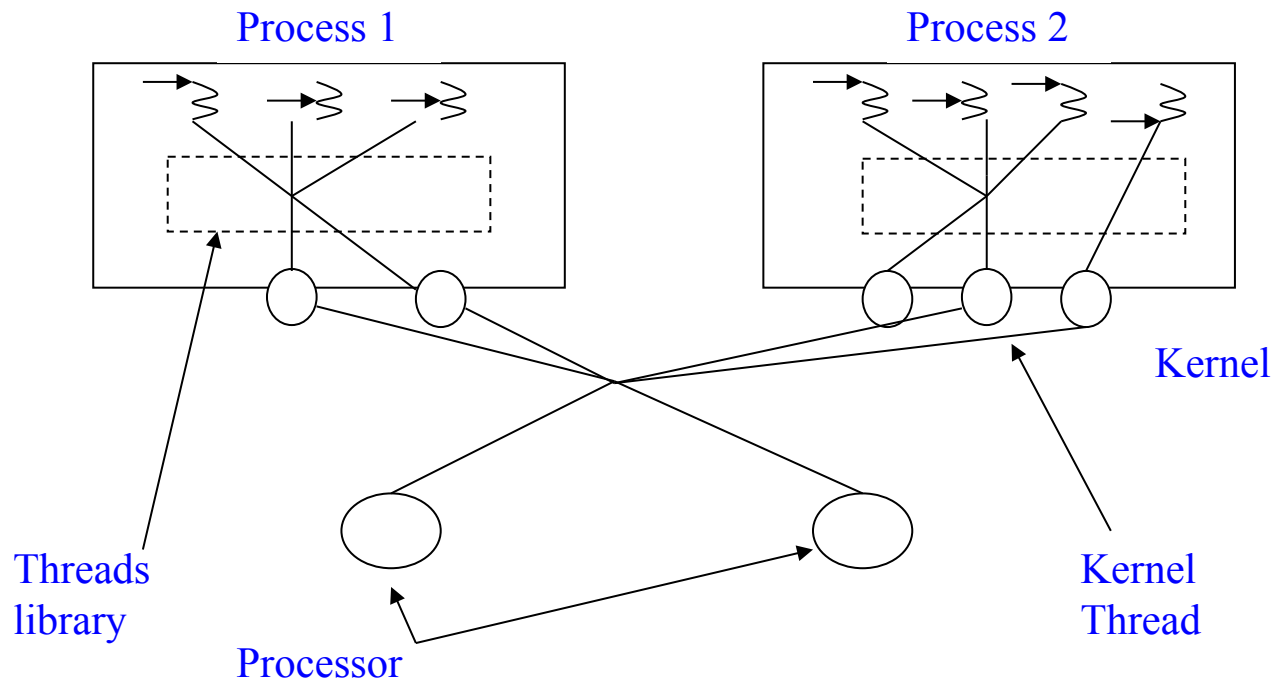
- ★ Each user thread is mapped into a kernel thread
 - allows more concurrency; allowing another thread to run when a thread is blocked.
 - Allows multiple threads running in parallel on multiprocessors.
- ★ Examples: Windows OS, Linux



Mapping (cont.)

Many-to-many Model:

- * Allows M user-level threads to be mapped to N kernel threads.
- * Allows the OS to create a sufficient number of kernel threads.
- * Many user threads can be created as necessary, and their corresponding kernel threads can run in parallel on multiprocessors.
- * Solaris 2, IRIX, HP-UX, Tru64 UNIX, Windows NT/2000 with the *ThreadFiber* package



Thread Libraries

- ★ Programmers use API of a thread library to create and manage threads
- ★ Thread library can be implemented as
 - User level threads
 - Kernel level threads
- ★ Three main thread libraries:
 - POSIX Pthreads → may be user or kernel threads
 - Win32 → kernel threads.
 - Java → implemented using the thread library of the host system
 - ★ In Windows: it uses Win32
 - ★ In Linux: use Pthreads
- ★ Read Textbook for example thread programs

Thread Libraries (cont.)

Pthreads

- ★ A POSIX standard (IEEE 1003.1c) API for thread creation and synchronization.
- ★ Primary goal: To allow multithreaded programs to be portable across multiple OS platforms.
- ★ API specifies behavior of the thread library, implementation is up to development of the library
 - can be kernel or user-level threads.
 - Common in UNIX operating systems.
- ★ POSIX: each thread maintains processor registers, stack, and signal mask
 - other resources must be globally accessible to all threads in the process.

Threading Issues

Semantics of fork() and exec() system calls:

- ★ If one thread invokes the exec() system call, the program specified in the parameter to exec() will replace the entire process, including all threads and LWPs.
- ★ If one thread in a program calls fork(), does the new process duplicate all threads or is the new process single-threaded (only the thread that invoked the fork)?
 - Duplicate only the thread that invoked the fork system call
 - ★ Appropriate if exec is called immediately.
 - Duplicate all threads
 - ★ if the new process does not call exec.

Threading Issues (cont.)

Thread cancellation:

- ★ A thread (target thread) is terminated before it has completed.
- ★ Two different scenarios:
 - Asynchronous cancellation – One thread immediately terminates the target thread.
 - Deferred cancellation – the target thread periodically checks if it should terminate.
- ★ Problems with cancellation:
 - If resources have been allocated to a cancelled thread → OS reclaims system resources, (usually) not all resources;
 - ★ asynchronous cancellation may not reclaim all resources.
 - A thread was cancelled while in the middle of updating data shared with other threads.

Threading Issues (cont.)

Signal handling:

- ★ A signal is used in Unix to notify a process that a particular event has occurred.
- ★ Synchronous signals: delivered to the same process that performs the operation causing the signal;
 - Operations can be divide by zero, illegal memory access, etc.
- ★ Asynchronous signals: generated by an event external to a running process; <control> <C>, timer expired.
- ★ All signals follow the same pattern:
 - Signal generation → occurrence of a particular event.
 - Delivering the signal to a process.
 - Signal handling.
- ★ Two possible signal handlers:
 - By a default signal handler → every signal has a default handler which is run by the kernel.
 - By a user-defined signal handler → override the default.

Threading Issues (cont.)

Signal handling (cont.):

- ★ When a process has more than one threads, where should the signal be delivered? Options:
 - Deliver the signal to the thread to which the signal applies → for synchronous signals.
 - Deliver the signal to every thread in the process → for some asynchronous signals such as <control><C>.
 - Deliver the signal to certain threads in the process → options in some versions of UNIX.
 - Assign a specific thread to receive all signals for the process → implemented in Solaris 2.
 - POSIX: signals are sent to process, and each thread has a signal mask to allow the thread to disable signals of a particular type
 - ★ A signal will be delivered to all threads in the process that are not masking signals of that type.

Threading Issues (cont.)

Thread pools

- ★ In a multithreaded web server, when the server receives a request, a separate thread is created; Disadvantages:
 - Time consuming for creating and destroying a thread for each request.
 - No limit (bound) on the number of threads created
 - ★ May exhaust system resources.
 - One solution: use thread pools
 - ★ Create a number of threads at process startup, and place them into a pool waiting for work.
- ★ The benefits of thread pools are:
 - Faster to service a request.
 - Limits the number of threads that exist at any one point.
- ★ The number of threads in the pool can be set:
 - Heuristically
 - ★ Based on system resources, such as CPU, memory or based on number of requests.
 - Dynamically
 - ★ Adjusted based on usage patterns.

Threading Issues (cont.)

Thread specific data:

- ★ By default, threads within a process share data of the process.
- ★ For some applications, each thread might need its own copy of certain data
 - They need thread-specific data.
- ★ Win32, Pthreads, Java support thread-specific data.

OS Examples

Windows Threads

- ★ Implement Windows API
- ★ Use one-to-one mapping
- ★ A process contains one or more threads
- ★ Each thread contains
 - a thread id
 - register set
 - separate user and kernel stacks
 - private data storage area
- ★ Context of a thread: register set, stacks, and private storage area

OS Examples

Linux Threads

- ★ Linux does not differentiate processes and threads
 - It refers to them as *tasks* rather than *threads* or *processes*
- ★ Linux provides both `fork ()` and `clone ()` system calls
 - `Fork ()` is used to create a *traditional* process
 - `Clone ()` is used to create a thread
 - ★ Can set flags to determine how much resource sharing between the parent and child tasks
 - ★ `Clone()` allows a child task to share the address space of the parent task (process).
- ★ Kernel version 2.6: one-to-one thread mapping.