

**COMMONWEALTH OF AUSTRALIA**  
**Copyright Regulation 1969**

**WARNING**

This material has been copied and communicated to you by or on behalf  
of **Curtin University of Technology** pursuant to Part VB of the  
*Copyright Act 1968* (**the Act**)

The material in this communication may be subject to copyright under the  
Act. Any further copying or communication of this material by you  
may be the subject of copyright protection under the Act.

Do not remove this notice

# **Operating Systems**

## **COMP2006**

### **Process Synchronization**

#### **Lecture 3**

# Process Synchronization

## References:

Silberschatz, Galvin, and Gagne, *Operating System Concepts*, Chapter 5.

## Topics:

- ★ Concurrent processes.
- ★ The critical section problem.
- ★ Solutions for critical-section problems using semaphores, etc.
- ★ Classical problems of synchronization.
- ★ Monitors

# Background

- ★ A *cooperating* process can affect or be affected by the other processes executing in the system.
  - an *independent* process cannot affect or be affected by other processes.
- ★ Cooperating processes may:
  - Share a logical address space (i.e., both code and data)
  - Share data through files
- ★ Concurrent access to shared data may result in data inconsistency.
- ★ Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes.

# Example: The producer/consumer problem

**Bounded buffers** (Circular list).

// in = out when the buffer is empty.

in: next free buffer;

out: first full buffer;

**Shared data**

**Type** *item* = ... ;

**var**        *buffer*: **array**[0..*n*-1] of *item*;  
             *nextp*, *nextc*: *item*;  
             *in*, *out*: 0 .. *n*-1;  
             *counter*: 0 .. *n*;  
             *in*, *out*, *counter* = 0;

**Producer process**

**repeat**

    produce an item in *nextp*;

    ...

*/\* busy waiting \*/*

**while** *counter* = *n* **do** *no-op*;

*buffer*[*in*] = *nextp*;

*in* = (*in*+1) **mod** *n*;

*counter* := *counter* + 1;

**until** *false*;

**Consumer process**

**repeat**

*/\* busy waiting \*/*

**while** *counter* = 0 **do** *no-op*;

*nextc* = *buffer*[*out*];

*out* = (*out*+1) **mod** *n*;

*counter* = *counter* - 1;

    consume item in *nextc*;

    ...

**until** *false*;

## Example (cont.)

- \* Each of the following statements must be executed *atomically*  $\rightarrow$  one uninterruptible unit.  
 $counter = counter + 1$ ; and  $counter = counter - 1$ ;
- \* *Atomic* operation means an operation that *completes entirely without interruption*.

**Example:** Assume initially  $counter = 5$ ;

producer executes  $counter = counter + 1$ ; and

consumer executes  $counter = counter - 1$ ;

- \* The correct result for  $counter$  should be 5
  - if the producer and consumer executes separately (not concurrently); or
  - If each of the instruction is executed atomically
  - Otherwise,  $counter$  can be 4, 5, or 6.
- \* Statements  $counter = counter + 1$  and  $counter = counter - 1$  may be implemented in machine language as:

$Register_1 = counter$ ;	$Register_1 = counter$ ;
$Register_1 = Register_1 + 1$ ;	$Register_1 = Register_1 - 1$ ;
$counter = Register_1$ ;	$counter = Register_1$ ;

## Example (cont.)

One example of interleaving of machine instruction:

$T_0$ : <i>producer</i> <b>execute</b>	$register_1 = counter$	$\{register_1 = 5\}$
$T_1$ : <i>producer</i> <b>execute</b>	$register_1 = register_1 + 1$	$\{register_1 = 6\}$
$T_2$ : <i>consumer</i> <b>execute</b>	$register_2 = counter$	$\{register_2 = 5\}$
$T_3$ : <i>consumer</i> <b>execute</b>	$register_2 = register_2 - 1$	$\{register_2 = 4\}$
$T_4$ : <i>producer</i> <b>execute</b>	$counter = register_1$	$\{counter = 6\}$
$T_5$ : <i>consumer</i> <b>execute</b>	$counter = register_2$	$\{counter = 4\}$

→ We get  $counter = 4$  (incorrect) because we allowed both processes to manipulate *counter* concurrently.

# Race condition

- ★ *Race condition* is a situation where several processes access and manipulate the same data concurrently.
  - The outcome of the execution depends on particular order in which the access takes place.
- ★ In order to prevent race condition on *counter*, we need to ensure that only one process at a time can be manipulating *counter*
  - We need some form of *process synchronization*.



# The critical section problem

- ★  $n$  processes  $\{P_0, P_1, \dots, P_{n-1}\}$  all competing to use some shared data.
- ★ Each process has a code segment, called *critical section*, in which the shared data is accessed.
- ★ The *critical section problem* is to make sure that when one process is executing in its critical section, no other process is allowed to execute in its critical section.

Structure of process  $P_i$

**Repeat**

*Entry section.*

        Critical section.

*Exit section.*

        Remainder section.

**Until** *false*;

# Solution to the critical-section problem

## Must satisfy these three requirements:

1. **Mutual Exclusion:** If a process is in the critical section then no other processes can be executing in their critical sections  
or No processes may be simultaneously inside their critical sections.
2. **Progress:** If no process is in its critical section and there exists some processes that wish to enter their critical sections, then only processes not in their remainder section can participate in the decision as to which one enters its critical section next (cannot postpone indefinitely)  
or No process running outside its critical section may block other processes from entering their critical sections.
3. **Bounded Waiting:** There must exist a bound on the number of times that other processes can enter their critical section after a process has made a request to enter its critical section and before that request is granted  
or No process should have to wait forever to enter its critical section.

## Assumption:

- \* Each process executes at non-zero speed.
- \* No assumption concerning *relative* speed of the  $n$  processes or the number of CPUs

## Solution (cont.)

### Simplest solution:

Each process *disables* all interrupts just *after entering* its critical section and *re-enables* them just *before leaving* it.

- Not wise, because enabling/disabling interrupt is a privileged instruction.

### Solution in kernel mode:

- ★ **Preemptive kernel:** a process can be preempted while running in the kernel mode
  - Otherwise, the kernel is **nonpreemptive kernel**: allows the process to run until it exits kernel mode, blocks, or voluntarily yields CPU.
- ★ Nonpreemptive kernel is free from race conditions on kernel data structure
  - However preemptive kernel is more responsive and suitable for real time system.

# Solution for two processes, $P_0$ and $P_1$

## Software-based

### ALGORITHM 1

**var** *turn*: (0 .. 1); // Initially *turn* = 0; *turn* = *i* means  $P_i$  can enter its CS

Process  $P_i$

**repeat**

**while** *turn*  $\neq i$  **do** *no-op*;

        Critical Section

*turn* = *j*;

        Remainder Section

**until** *false*;

Process  $P_j$

**repeat**

**while** *turn*  $\neq j$  **do** *no-op*;

        Critical Section

*turn* = *i*;

        Remainder Section

**until** *false*;

- \* Satisfies mutual exclusion, but not progress requirement.
  - If *turn* = 0,  $P_1$  cannot enter its CS even though  $P_0$  is in its RS.
  - Taking turn is not good when one process is slower than other.
- \* *Busy waiting*: continuously testing a variable waiting for some value to appear
  - not good since it wastes CPU time

# Solution for two processes (cont. )

## Software-based

### ALGORITHM 2

// Initially  $flag[0] = flag[1] = false$ ;  $flag[i] = true$  means  $P_i$  wants to enter its CS

**var** *flag*: **array** [0 .. 1] **of** *boolean*;

Process  $P_i$

**repeat**

$flag[i] = true$ ;

**while**  $flag[j]$  **do** *no-op*;

        Critical Section;

$flag[i] = false$ ;

        Remainder Section;

**until** *false*;

Process  $P_j$

**repeat**

$flag[j] = true$ ;

**while**  $flag[i]$  **do** *no-op*;

        Critical Section;

$flag[j] = false$ ;

        Remainder Section;

**until** *false*;

\* Satisfy mutual exclusion, but violates the progress requirement:

$T_0$ :  $P_0$  sets  $flag[0] = true$ .

$T_1$ :  $P_1$  sets  $flag[1] = true$ .

→  $P_0$  and  $P_1$  are looping in their respective **while**.

# Solution for two processes (cont. )

## Software-based

// **Peterson's solution:** Combine shared variables of Algorithms 1 and 2

Process  $P_i$

**repeat**

*flag[i] = true;*

*turn = j;*

**while** (*flag[j] and turn = j*) **do no-op;**

critical section

*flag[i] = false;*

remainder section

**until** *false;*

Process  $P_j$

**repeat**

*flag[j] = true;*

*turn = i;*

**while** (*flag[i] and turn = i*) **do no-op;**

critical section

*flag[j] = false;*

remainder section

**until** *false;*

- ★ Solves the critical-section problem for two processes.
  - It meets all the three requirements
- ★ Proof: need to show that:
  - Mutual exclusion is preserved.
  - The progress requirement is satisfied.
  - The bounded waiting time requirement is met.
- ★ For detailed proof, read the textbook.

# Bakery Algorithm

- ★ The solution to the critical section problem for  $n$  processes by Leslie Lamport
- ★ Before entering its critical section, each process receives a number.
  - The holder of the smallest number enters the critical section.
  - If processes  $P_i$  and  $P_j$  receive the same number, if  $i < j$ , then  $P_i$  is served first; else  $P_j$  is served first.
- ★ Notation (**ticket#**, process **id#**)
  - $(a, b) < (c, d)$  if  $a < c$  or if  $a = c$  and  $b < d$ .
  - $\max(a_0, \dots, a_{n-1})$  is a number  $k$  such that  $k \geq a_i$  for  $i = 0, \dots, n-1$ .
- ★ shared data
  - **var** *choosing*: **array**[0.. $n-1$ ] of *boolean*;
  - *number*: **array** [0 ..  $n-1$ ] of *integer*;
- ★ data structures are initialised to *false* and 0, respectively

# Bakery Algorithm (cont.)

Process  $P_i$

repeat

$choosing[i] = true;$

$number[i] = \max(number[0], number[1], \dots, number[n-1]) + 1;$

$choosing[i] = false;$

for  $j = 0$  to  $n-1$  do

begin

while  $choosing[j]$  do no-op;

while  $number[j] \neq 0$  and  $(number[j], j) < (number[i], i)$  do no-op;

end;

critical section

$number[i] = 0;$

remainder section

until *false*



# Synchronization Hardware

- ★ There is no guarantee that the software-based solution will work correctly in all computer architectures.
- ★ The simple solution to critical section problem: disable interrupt while a shared variable is being modified.
  - This solution is not feasible in multiprocessor. Why?
- ★ Use special hardware instructions such as *Test-and-Set* and *Swap*.
  - *Test-and-set* or *Swap* is an atomic instruction: it can not be interrupted until it completes its execution

// Test and set the content of a word  
*atomically*

```
function Test-and-Set (var boolean: target)  
begin  
    Test-and-Set = target;  
    target = true;  
end;
```

// Swapping instruction is done  
*atomically*

```
procedure Swap (var boolean: a, b)  
var boolean: temp;  
begin  
    temp = a;  
    a = b;  
    b = temp;  
end;
```

# How to use them?

## Mutual Exclusion with *Test-and-Set*

var *boolean: lock*; *lock* is a shared variable, initially set to *false*.

```
Repeat // Process Pi
    while Test-and-Set (lock) do no-op;
        Critical Section
    lock = false;
        Remainder Section
until false;
```

```
Repeat // Process Pj
    while Test-and-Set (lock) do no-op;
        Critical Section
    lock = false;
        Remainder Section
until false;
```

## Mutual Exclusion with *Swap*

```
Repeat // Process Pi
    key = true;
    repeat
        Swap (lock, key);
    until key = false;
        Critical section
    lock = false;
        Remainder section
until false;
```

```
Repeat // Process Pj
    key = true;
    repeat
        Swap (lock, key);
    until key = false;
        Critical section
    lock = false;
        Remainder section
until false;
```

★ Both do not satisfy the bounded waiting requirement.

# Correct solution with Test-and-set

Shared data: **var** *waiting*: **array**[0..*n*-1] **of** *boolean*; *lock*: *boolean*; //All initialized to *false*

**Process  $P_i$**

**var** *j*: 0..*n*-1; *key*: *boolean*;

**repeat**

*waiting*[*i*] = *true*;

*key* = *true*;

**while** *waiting* [*i*] **and** *key* **do** // enter CS if either *waiting*[*i*] or *key* is *false*

*key* = Test-and-Set (*lock*); // *key* is *false* if *lock* is *false*

*waiting*[*i*] = *false*;

*Critical Section*

*j* = *i*+1 **mod** *n*;

**while** (*j* ≠ *i*) **and not** *waiting*[*j*] **do** // check if any  $P_j$  is waiting for CS

*j* = *j*+1 **mod** *n*

**if** *j* = *i* **then**

*lock* = *false*; // no other process is waiting for CS

**else**

*waiting*[*j*] = *false*; //  $P_j$  is waiting, let it enter CS next

*Remainder Section*

**until** *false*;

**Proof:** Read textbook.

# Mutex locks

- ★ The hardware solutions are complicated and inaccessible to application programmers
  - Use software tool, called mutual exclusion (mutex)
- ★ Mutex: acquire a lock before *entering* a critical section, and release the lock on *exit*
  - Calls to *acquire ()* or *release ()* must be atomic
  - Implement the functions with the hardware solution (e.g., test-and-set)
  - Busy waiting (also called *spinlock*) wastes CPU, but does not need context switch → it is good with short wait.
    - ★ Implemented on multiprocessor systems

## Repeat

*acquire lock.*

Critical section.

*release lock*

Remainder section.

**Until** *false*;

```
acquire () {  
    while (!available); // busy wait  
    available = false  
}  
  
release () {  
    available = true  
}
```

# Semaphores

- ★ Semaphore  $S$  is an *integer variable* that can only be accessed via two indivisible (atomic) operations.

*wait* ( $S$ ):     **while**  $S \leq 0$  **do** *no-op*; // busy wait  
                   $S = S - 1$ ;

*signal* ( $S$ ):  $S = S + 1$ ;

- ★ *wait* ( $S$ ) and *signal* ( $S$ ) are respectively equivalent to  $P(S)$  and  $V(S)$ 
  - *Proberen* means to test; *Verhogen* means to increment
- ★ Can be used to solve the  $n$ -process critical section problem.
- ★ Note, the definition of *wait* ( $S$ ) and *signal* ( $S$ ) uses *busy waiting*.
  - Also called *spinlock*; that is, the process *spins* waiting for the *lock* to be false
  - Disadvantage? Advantage?

## Semaphores (cont.)

Shared variables: **var** *mutex: semaphore*;

*/\* initially mutex = 1 \*/*

Process  $P_i$ :

**Repeat**

*wait (mutex);*

        Critical Section

*signal (mutex);*

        Remainder Section

**until** *false*;

Process  $P_j$ :

**Repeat**

*wait (mutex);*

        Critical Section

*signal (mutex);*

        Remainder Section

**until** *false*;

Process  $P_k$ :

**Repeat**

*wait (mutex);*

        Critical Section

*signal (mutex);*

        Remainder Section

**until** *false*;

# Semaphore Implementation

// Define a semaphore as a record

**type semaphore = record**

*value*: integer;

*list*: **list of process**;

**end;**

*wait (S):* *S.value* = *S.value* - 1;

**if** *S.value* < 0 **then**

**begin**

            add process to *S.list*;

*block* ();

**end;**

*signal (S):* *S.value* = *S.value* + 1;

**if** *S.value* ≤ 0 **then**

**begin**

            remove a process *P* from *S.list*;

*wakeup* (*P*);

**end;**

- *block* () suspends the process that invokes it
  - The process is **put into a waiting queue** associated with the semaphore; i.e., *S.list*
  - The process state is switched from **running** into **waiting** state.
  - CPU scheduler selects another process from **ready queue**.
- *wakeup* (*P*) resumes the execution of a blocked process *P* in *S.list*
  - *P* is placed in the **ready queue**.
  - The process state is switched from **waiting** into **ready** state.
- Each *wait* and *signal* must be executed atomically.
  - In a uni-processor, inhibit interrupts during *wait* and *signal*.
    - OK because it is done by the system
  - In a multi-processor, use other method such as spinlock.

# Semaphore as General Synchronization Tool

- ★ Consider two processes that require  $P_j$  to execute code  $B$  only after code  $A$  has been executed by  $P_i$ .
- ★ Use semaphore  $flag$  initialized to 0.

## Code:

$P_i$

.

.

$A$

$signal(flag)$

$P_j$

.

.

$wait(flag)$

$B$

- ★ *Counting* semaphore – integer value can range over an unrestricted domain
- ★ *Binary* semaphore – integer value can range only between 0 and 1
  - can be simpler to implement.



# Deadlock and Starvation

- ★ Deadlock – two or more processes are *waiting indefinitely* for an event that can be caused by only one of the waiting processes.
- ★ Let  $S$  and  $Q$  be two semaphores initialized to 1.

$P_0$

*wait* ( $S$ );

*wait* ( $Q$ );

.

.

*signal* ( $S$ );

*signal*( $Q$ );

$P_1$

*wait* ( $Q$ );

*wait* ( $S$ );

.

.

*signal*( $Q$ );

*signal* ( $S$ );

- ★ Other problem: Starvation – *indefinite blocking*.
  - A process may never be removed from the semaphore queue in which it is suspended.
  - What happens if  $S.list$  is implemented as a last-in first out (LIFO)?

# Priority Inversion

- ★ **Scenario 1:** Consider a shared kernel data  $X$  that is being accessed by a process  $A$ . What happens if a higher priority process  $B$  wants to access  $X$ ?
  - Kernel data is usually protected by a lock, and thus  $B$  must wait for  $A$  to finish with the data
- ★ **Scenario 2:** Consider a shared resource  $R$  and three processes,  $A$ ,  $B$  and  $C$ , where  $C$  and  $A$  have the highest and lowest priority respectively.
  - Assume  $R$  is being used by  $A$ , and  $C$  wants to use  $R$ 
    - ★  $C$  must wait until  $A$  finishes with  $R$
  - What happens if  $B$  preempts  $A$ ? ( $B$  has higher priority than  $A$ , but lower than  $C$ )
    - ★ Process  $C$  must wait for process  $B$  that has a lower priority before it can access resource  $R$ .
    - ★ This is called *priority inversion* problem
    - ★ Priority inversion occurs only in a system with more than two priorities.

**Solution:** *priority inheritance protocol* – all processes that are accessing resources needed by a higher priority process inherits the higher priority until they finish using the resources

- ★ Process  $A$  would temporarily have the priority of process  $C$ , and thus  $B$  cannot preempt  $A$
- ★ When  $A$  finish using  $R$ , it goes back to its own priority and  $C$  will run because it has a higher priority than  $B$

# Bounded-Buffer Problem

Shared data

```
type item = ...  
var buffer = ...  
    mutex, full, empty: semaphore;  
    nextp, nextc: item;  
    full = 0; empty = n; mutex = 1;
```

Producer process

**repeat**

```
    ...  
    produce an item in nextp  
    ...  
    wait (empty);  
    wait (mutex);  
    ...  
    add nextp to buffer  
    ...  
    signal (mutex);  
    signal (full);
```

```
    ...  
until false;
```

Consumer process

**repeat**

```
    wait (full);  
    wait (mutex);  
    ...  
    remove item from buffer to nextc  
    ...  
    signal (mutex);  
    signal (empty);  
    ...  
    consume item in nextc
```

```
    ...  
until false;
```

# Readers-writers problem

- ★ Synchronisation problem involving reading and writing shared data objects.
  - Data can be accessed by more than one reader at the same time
  - However, when the data is being accessed by a writer, no other writer or readers can access it simultaneously.
- ★ Assume no reader will be kept waiting unless a writer has already obtained permission to use shared object.
  - reader has higher priority.
  - An alternative problem: writer has higher priority than reader.
  - Solutions to either problem may cause starvation

# Readers-writers problem (cont.)

Shared data

```
var mutex, wrt: semaphore;  
    readcount: integer;  
    readcount := 0; /* number of processes currently reading object */  
    mutex := 1; /* ensure mutual exclusion for readcount */  
    wrt := 1 /* for writers critical section and used by the first and last readers */
```

## Reader process

**repeat**

```
    P(mutex);  
        readcount = readcount + 1;  
        if readcount = 1 then  
            P(wrt);  
    V(mutex);  
    ...  
    reading is performed;  
    ...  
    P(mutex);  
        readcount = readcount - 1;  
        if readcount = 0 then  
            V(wrt);  
    V(mutex);
```

**until false;**

## Writer process

**repeat**

```
    P(wrt);  
    ...  
    writing is performed  
    ...  
    V(wrt);
```

**until false;**

# Readers-writers problem


- ★ Variable *mutex* is used to ensure mutual exclusion on variable *readcount*
- ★ If writer is in its critical section and  $n$  readers are waiting, then:
  - one reader process is queued on *wrt*, and
  - $n-1$  reader processes are queued on *mutex*.
- ★ Reader-writer problem and its solution has been used to provide reader-writer locks on some systems, e.g., Linux.
  - When a process wants to only read shared data, it gets the lock as reader.
    - ★ More than one processes can read at the same time
  - When a process wants to write shared data, it gets the lock as writer
    - ★ Only one process at a time can get the write lock.

# Dining-philosophers problem by Dijkstra

- ★ Five philosophers sit around a circular table
  - Each philosopher has a bowl of spaghetti
  - There are five *chopsticks*, each is placed between each pair of bowls
  - Each philosopher needs *two chopsticks* to eat the spaghetti
- ★ Each philosopher can be either *thinking* or *eating*
  - When *thinking*, the philosopher doesn't interact with the others
  - From time to time each philosopher gets *hungry*
    - ★ When getting *hungry*, a philosopher tries to *grab* two chopsticks close to her (i.e., one on her left and the other on her right) from the table
      - A philosopher can pick up one chopstick at a time
      - A philosopher is not allowed to seize chopsticks from others' hands! Not polite!
  - A philosopher can eat once she has two chopsticks
    - ★ Once finish eating, she *releases* both chopsticks, and start thinking again

**Problem:** synchronize the philosophers so that each can think and eat with no deadlock and no philosopher can be starved to death!

# Dining-philosophers problem (cont. )



Shared data

**var** *chopstick*: **array**[0..4] **of** *semaphore*; // initially = 1

Philosopher *i*

**repeat**

*P*(*chopstick*[*i*]); // pick left chopstick

*P*(*chopstick*[(*i*+1) **mod** 5]); // pick right chopstick

...

*eat* // you have two chopsticks

...

*V*(*chopstick*[*i*]); // release left chopstick

*V*(*chopstick*[(*i*+1) **mod** 5]); // release right chopstick

...

*think* // you have no chopsticks

...

**until** *false*;

**In the solution:** No two neighbours can eat simultaneously; may have a deadlock (if all five grab left chopsticks).

## **Other Solutions:**

- \* Allow at most four philosophers to sit at the table.
- \* Pickup chopsticks only if both are available; picking both chopsticks is a critical section
- \* *Odd* philosopher picks left chopstick first; *even* philosopher grabs right chopstick first.



# Semaphore Limitation

- ✱ Incorrect use of semaphore can still result in timing errors → hard to detect.
- ✱ Omitting  $P$ , or  $V$ , or both may create deadlock or no mutual exclusion.

- ✱ **Example 1:** several processes may execute in their critical sections simultaneously:

$V(\text{mutex});$

...

critical section

...

$P(\text{mutex});$

- ✱ **Example 2:** deadlock may occur:

$P(\text{mutex});$

...

critical section

...

$P(\text{mutex});$

# Monitors

- ★ The Monitor is a high-level synchronization construct that allows the safe sharing of an abstract data type among concurrent processes.
- ★ A monitor type contains a set of **defined functions** and a set of **shared variables** that can be accessed only by one of the functions
  - The monitor ensures that only one process at a time is active in the monitor

**type** *monitor-name* = **monitor**

variable declarations

**procedure entry**  $P_1$  (...);

**begin ... end;**

**procedure entry**  $P_2$  (...);

**begin ... end;**

**procedure entry**  $P_n$  (...);

**begin ... end;**

**begin**

initialisation code

**end.**




Fig. 5.16 (textbook)

# Monitors (cont.)

- \* To allow a process to wait within the monitor, a condition variable must be declared as:

**var**  $x, y$ : *condition*

- \* Condition variable can only be used with the operations *wait* and *signal*.
- \* The operation:

$x.wait$ ;

means that the process invoking this operation is suspended until another process invokes

$x.signal$

- \* the  $x.signal$  operation resumes exactly one suspended process. If no process is suspended, then the signal operation has no effect → different from semaphore.




Fig. 5.17 (textbook)

## Monitors (cont. )

- ★ What happens if there is a suspended process  $Q$  waiting for condition  $x$  when  $x.signal()$  is called by process  $P$ ?
  - Monitor must ensure that only one of the two processes can be active at a time.
    - ★ If  $Q$  is activated,  $P$  must wait
- ★ There are two possibilities:
  - Signal and wait:  $P$  either waits until  $Q$  leaves the monitor or waits for another condition
  - Signal and continue:  $Q$  either waits until  $P$  leaves the monitor or waits for another condition

# Dining Philosophers example

**type** *dining-philosophers* = **monitor**

**var** *state*: **array** [0 .. 4] **of** (*thinking*, *hungry*, *eating*);

**var** *self*: **array** [0 .. 4] **of** *condition*;

**procedure entry** *pickup* (*i*: 0 .. 4);

**begin**

*state*[*i*] = *hungry*;

*test* (*i*);

**if** *state*[*i*] ≠ *eating* **then** *self*[*i*].*wait*;

**end**;

**procedure entry** *putdown* (*i*:0..4);

**begin**

*state*[*i*] = *thinking*;

*test* (*i*+4 **mod** 5);

*test* (*i*+1 **mod** 5);

**end**;

**// initialization**

**begin**

**for** *i* = 0 to 4 **do**

*state*[*i*] = *thinking*;

**end**.

**procedure** *test* (*k*: 0 .. 4);

**begin**

    //two neighbors are not eating

**if** *state*[*k*+4 **mod** 5] ≠ *eating*

**and** *state* [*k*] = *hungry*

**and** *state* [*k*+1 **mod** 5] ≠ *eating* **then**

**begin** //pick up chopsticks if both are available

*state*[*k*] = *eating*;

*self*[*k*].*signal*;

**end**;

**end**;

★ Philosopher *i* must invoke:

*dp*.*pickup*(*i*);

    ...

*dp*.*putdown*(*i*);

★ Still possibility of starvation.

# Semaphores for Monitor

- ★ Use a *mutex* variable (initialized to one) for each monitor
  - A process executes *wait (mutex)* before entering the monitor
  - Use *signal (mutex)* when exiting the monitor
  - See Section 5.8.3 for details!